

We are now Refinitiv, formerly the Financial and Risk business of Thomson Reuters. We've set a bold course for the future – both ours and yours – and are introducing our new brand to the world.

As our brand migration will be gradual, you will see traces of our past through documentation, videos, and digital platforms.

Thank you for joining us on our brand journey.



REFINITIVTM



Transport API C Edition V3.2.x

OPEN SOURCE PERFORMANCE TOOLS GUIDE C EDITION



© Thomson Reuters 2016 - 2018. All rights reserved.

Thomson Reuters, by publishing this document, does not guarantee that any information contained herein is and will remain accurate or that use of the information will ensure correct and faultless operation of the relevant service or equipment. Thomson Reuters, its agents and employees, shall not be held liable to or through any user for any loss or damage whatsoever resulting from reliance on the information contained herein.

This document contains information proprietary to Thomson Reuters and may not be reproduced, disclosed, or used in whole or part without the express written permission of Thomson Reuters.

Any Software, including but not limited to, the code, screen, structure, sequence, and organization thereof, and Documentation are protected by national copyright laws and international treaty provisions. This manual is subject to U.S. and other national export regulations.

Nothing in this document is intended, nor does it, alter the legal obligations, responsibilities or relationship between yourself and Thomson Reuters as set out in the contract existing between us.

Contents

Chapter 1	Introduction	1
1.1	About this Manual	1
1.2	Audience	1
1.3	Programming Language.....	1
1.4	Acronyms and Abbreviations	1
1.5	References	2
1.6	Document Conventions.....	2
1.6.1	<i>Typographic</i>	2
1.6.2	<i>Diagrams</i>	3
Chapter 2	Open Source Performance Tool Suite Overview.....	4
2.1	Overview	4
2.2	Transport API Performance Tool Suite	4
2.3	Package Contents.....	6
2.3.1	<i>Building</i>	6
2.3.2	<i>Running</i>	6
2.4	What Is Measured and Reported	7
2.4.1	<i>Latency</i>	7
2.4.2	<i>Throughput and Payload</i>	7
2.4.3	<i>Image Retrieval Time</i>	7
2.4.4	<i>CPU & Memory Usage</i>	7
2.5	Recorded Results and Output.....	8
2.5.1	<i>Summary File</i>	8
2.5.2	<i>Statistics File</i>	8
2.5.3	<i>Latency File</i>	8
Chapter 3	Latency Measurement Details.....	9
3.1	Time-slicing	9
3.2	Latency.....	10
Chapter 4	upacConsPerf	12
4.1	Overview	12
4.2	Threading and Scaling	12
4.2.1	<i>Consumer Lifecycle</i>	12
4.2.2	<i>Diagram</i>	14
4.3	Latency Measurement.....	15
4.3.1	<i>Consumer Latency</i>	15
4.3.2	<i>Posting Latency</i>	15
4.4	upacConsPerf Configuration Options.....	16
4.5	Input	18
4.6	Output	18
4.6.1	<i>upacConsPerf Summary File Sample</i>	19
4.6.2	<i>upacConsPerf Statistics File Sample</i>	20
4.6.3	<i>upacConsPerf Latency File Sample</i>	21
4.6.4	<i>upacConsPerf Console Output Sample</i>	21
Chapter 5	upacProvPerf	22
5.1	Overview	22
5.2	Threading and Scaling	22
5.3	Provider Lifecycle.....	22

5.4	Latency Measurement.....	24
5.5	upacProvPerf Configuration Options.....	24
5.6	Input Files.....	26
5.7	Output	27
5.7.1	<i>upacProvPerf Summary File Sample</i>	27
5.7.2	<i>upacProvPerf Statistics File Sample</i>	28
5.7.3	<i>upacProvPerf Console Output Sample</i>	28
Chapter 6	upacNIProvPerf	29
6.1	Overview	29
6.2	Threading and Scaling	29
6.3	Non-Interactive Provider Lifecycle	29
6.4	Latency Measurement.....	31
6.5	upacNIProvPerf Configuration Options.....	31
6.6	Input Files.....	33
6.7	Output	33
6.7.1	<i>upacNIProvPerf Summary File Sample</i>	34
6.7.2	<i>upacNIProvPerf Statistics File Sample</i>	35
6.7.3	<i>upacNIProvPerf Console Output Sample</i>	35
Chapter 7	upacTransportPerf	36
7.1	Overview	36
7.2	Threading and Scaling	36
7.3	upacTransportPerf Life Cycle.....	36
7.4	Message Payload.....	38
7.5	Latency Measurement.....	38
7.6	upacTransportPerf Configuration Options.....	38
7.7	Input	41
7.8	Output	41
7.8.1	<i>upacTransportPerf Summary File Sample</i>	41
7.8.2	<i>upacTransportPerf Statistics File Sample</i>	42
7.8.3	<i>upacTransportPerf Console Output Sample</i>	43
Chapter 8	Performance Measurement Scenarios	44
8.1	Interactive Provider to Consumer, Through TREP.....	44
8.2	Interactive Provider to Consumer, Direct Connect.....	45
8.3	Non-Interactive Provider to Consumer, Through TREP	46
8.4	Consumer Posting on the TREP	47
8.5	Transport Performance, Direct Connect with TCP	48
8.6	Transport Performance, Direct Connect with TCP, Reflection	49
8.7	Transport Performance, Direct Connect with Multicast	50
8.8	Transport Performance, Direct Connect with Shared Memory	51
Chapter 9	Input File Details	52
9.1	Message Content File and Format.....	52
9.1.1	<i>Encoding Fields</i>	52
9.1.2	<i>Sample Update Message</i>	53
9.1.3	<i>Sample MarketByOrder Data</i>	53
9.2	Item List File	56
9.2.1	<i>Item Attributes</i>	56
9.2.2	<i>Sample Item List File</i>	57
Chapter 10	Output File Details	58
10.1	Overview	58

10.2	Output Files and Their Descriptions	58
10.3	Latency File	59
10.4	File Import	59
Chapter 11	Performance Best Practices	61
11.1	Overview	61
11.2	Transport Best Practices	61
11.2.1	<i>rsslRead</i>	61
11.2.2	<i>rsslWrite, rsslFlush</i>	61
11.2.3	<i>Packing</i>	62
11.2.4	<i>High-water Mark</i>	62
11.2.5	<i>Direct Socket Write</i>	63
11.2.6	<i>Nagle's Algorithm</i>	63
11.2.7	<i>System Send and Receive Buffers</i>	63
11.2.8	<i>Transport API Buffering</i>	64
11.2.9	<i>Compression</i>	65
11.3	Encoder and Decoder Best Practices	65
11.3.1	<i>Single-Pass Encoding</i>	65
11.3.2	<i>Clear Functions</i>	65
11.4	Other Practices: CPU Binding	65
Appendix A	Troubleshooting	66
A.1	Can't Connect	66
A.2	Not Achieving Steady State	66
A.3	Consumer Tops Out but Not at 100% CPU	67
A.4	Initial Latencies Are High	67
A.5	Latency values Are Very High	67

List of Figures

Figure 1.	Running Performance Example and Host Notation	3
Figure 2.	Network Diagram Notation	3
Figure 3.	Three Connection Options for the OMM-based Performance Tools	4
Figure 4.	Transport API C Transport Perf	5
Figure 5.	Directory Structure of the Performance Tools	6
Figure 6.	Time Slicing Algorithm	9
Figure 7.	Refresh Publishing Algorithm	9
Figure 8.	Latency RICs within a Tick	10
Figure 9.	Timing Diagram for Latency Measurements	10
Figure 10.	upacConsPerf Lifecycle	12
Figure 11.	upacConsPerf Application Flow	14
Figure 12.	upacProvPerf Lifecycle	22
Figure 13.	upacProvPerf Application Flow	23
Figure 14.	upacNIProvPerf Lifecycle	29
Figure 15.	upacNIProvPerf Application Flow	30
Figure 16.	upacTransportPerf Lifecycle 1	36
Figure 17.	upacTransportPerf Application Flow	37
Figure 18.	Interactive Provider to Consumer on TREP	44
Figure 19.	Interactive Provider to Consumer, Direct Connect	45
Figure 20.	NIProv to Consumer on the TREP	46
Figure 21.	Consumer Posting to TREP	47
Figure 22.	Transport Performance, TCP Direct Connect	48
Figure 23.	Transport Performance, TCP Direct Connect with Reflection	49
Figure 24.	Transport Performance, Multicast Direct Connect	50
Figure 25.	TransportPerf, Shared Memory Direct Connect	51
Figure 26.	Sample Excel Graph from ConsStats1.csv	59
Figure 27.	Sample Excel Graph of Latencies Over a 15-second Steady State Interval from ConsLatency1.csv	60
Figure 28.	ADS rmds.cnf	67

List of Tables

Table 1:	Acronyms and Abbreviations	1
Table 2:	upacConsPerf Configuration Options.....	16
Table 3:	upacProvPerf Configuration Options.....	24
Table 4:	upacNIProvPerf Configuration Options	31
Table 5:	upacTransportPerf Configuration Options.....	38
Table 6:	Item Attributes	56
Table 7:	Performance Suite Applications and Associated Configuration Files	58

Chapter 1 Introduction

1.1 About this Manual

This guide introduces the Elektron Transport API C Edition of the performance suite. It presents an overview of how performance suite applications work with the Thomson Reuters Enterprise Platform (TREP), how the applications themselves work, and how application tests are run. It also provides an overview of the basic concepts of writing performant applications, as well as configuring both the applications and the Elektron Transport API for optimal performance.

Authors and contributors include Elektron Transport API architects and developers who encountered and resolved many of issues you might face. As such, this document is concise and addresses realistic scenarios and use cases.

This guide documents the general design and usage of the tools provided for measuring the performance. It describes how features of the API send and receive data with high throughput and low latency. This information applies both when the API connects directly to itself as well as when using intermediaries, such as TREP components like that Advanced Data Hub (ADH) and Advanced Distribution Server (ADS).

For simplicity, the Elektron Transport API is referred to throughout this manual as the Transport API.

1.2 Audience

This document is written to help programmers take advantage of Transport API features and achieve high throughput and low latency with their applications. The information detailed herein assumes that the reader is a user or a member of the programming staff involved in the design, code, and test phases for applications that will use the Transport API. It is assumed that you are familiar with the data types, operational characteristics, and user requirements of real-time data delivery networks, and that you have experience developing products using the C programming language in a networked environment. It is assumed that the reader has read the *Transport API C Edition Developer's Guide* to have a basic familiarity with the API Transport and the interaction models of OMM Consumers, OMM Interactive Providers, and OMM Non-Interactive Providers.

1.3 Programming Language

The Transport API C edition is written to the C language. All code samples in this document and all example applications provided with the product are written in C.

1.4 Acronyms and Abbreviations

ACRONYM	DEFINITION
ADH	Advanced Data Hub
ADS	Advanced Distribution Server
API	Application Programming Interface
CPU	Central Processing Unit
DMM	Domain Message Model
HTTP	Hypertext Transfer Protocol

Table 1: Acronyms and Abbreviations

ACRONYM	DEFINITION
HTTPS	Hypertext Transfer Protocol (Secure)
NIC	Network Interface Card
NIP	Non-Interactive Provider
OMM	Open Message Model
OS	Operating System
RAM	Random Access Memory
RDM	Reuters Domain Model
RFA	Robust Foundation API
RSSL	Reuters Source Sink Library
RWF	Reuters Wire Format
TREP	Thomson Reuters Enterprise Platform

Table 1: Acronyms and Abbreviations (Continued)

1.5 References

- *Transport API C Edition Developers Guide*
- *Transport API C Edition RDM Usage Guide*
- *Transport API C Edition Value Added Components Developers Guide*
- The [Thomson Reuters Professional Developer Community](#)

1.6 Document Conventions

1.6.1 Typographic

- C structures, methods, in-line code snippets, and types are shown in **orange, Lucida Console** font.
- Parameters, filenames, tools, utilities, and directories are shown in **Bold** font.
- Document titles and variable values are shown in *italics*.
- When initially introduced, concepts are shown in ***Bold, Italics***.
- Longer code examples (one or more lines of code) are shown in Lucida Console font against an orange background. Comments in the code are in green font. For example:

```
/* decode contents into the filter list structure */
if ((retVal = rsslDecodeFilterList(&decIter, &filterList)) >= RSSL_RET_SUCCESS)
{
    /* create single filter entry and reuse while decoding each entry */
    RsslFilterEntry filterEntry = RSSL_INIT_FILTER_ENTRY;
```

1.6.2 Diagrams

Diagrams that depict a component in a performance scenario use the following format. The grey box represents one physical machine, whereas blue or white boxes represent processes running on that machine.



Figure 1. Running Performance Example and Host Notation

Diagrams that depict the interaction between components on a network use the following notation:

	Feed Handler, Enterprise Platform server, or other application		Network of multiple servers
	UPA application		Point-to-point connection showing direction of primary data flow
	Shared memory segment		Point-to-point connection showing direction of client connecting to server
	Multicast network		Data from external source (e.g. consolidated network or exchange)
	Connection to Multicast network, no primary data flow direction		Connection to Multicast network showing direction of primary data flow

Figure 2. Network Diagram Notation

Chapter 2 Open Source Performance Tool Suite Overview

2.1 Overview

The general idea behind the Open Source Performance Tool Suite is to provide a consistent set of platform test applications that look and behave consistently across the Elektron APIs. The tool suite covers the various OMM-based API products and allows Thomson Reuters's internal and external clients to compare latency and throughput trade-offs of the various APIs and their differing functionality sets.

TREP also offers the tools **rmdsTestClient** and **sinkDrivenSrc** for performance testing, focusing on throughput, latency, and capacity of TREP components. The tool suite focuses on what can be done with each API and is meant to compliment other platform tools.

All tools in the suite are provided as buildable open-source and demonstrate best practice and coding for performance with their respective APIs. Future releases of API products will expand on these tests to include other areas of functionality (e.g., batch requesting, etc.). Clients can run these tools to determine performance results for their own environments, recreate Thomson Reuters-released performance numbers generated using these tools, and modify the open source to tune and tweak applications to best match their end-to-end needs.

These performance tools can generate reports comparing performance across all API products.

2.2 Transport API Performance Tool Suite

The Transport API C-based suite consists of an OMM consumer, OMM interactive provider, and OMM non-interactive provider. These applications showcase optimal OMM content consumption and providing within the TREP. Additionally, the Transport API provides a transport-only performance example which you can use to measure the performance of the Transport API transport handling opaque, non-OMM content. Source code is provided for all performance tool examples, so you can determine how functionality is coded and modify applications to suit your specific needs.

Because applications from the Elektron APIs are fully compatible and use similar methodologies, you can run them stand-alone within an API or mix them (e.g., a provider from Transport API and a consumer from RFA).¹

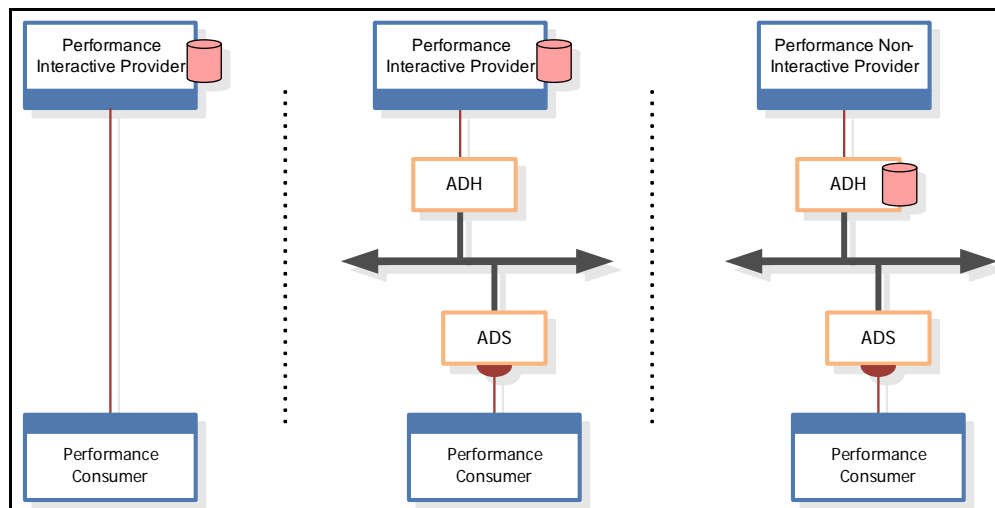


Figure 3. Three Connection Options for the OMM-based Performance Tools

1. Tools from the RFA C++ and RFA Java APIs must be obtained from their respective distribution packages.

In a typical OMM configuration, latency through the system is measured either one-way from a provider to consumer, or round-trip from a consumer, through the system, and back.² Latency information is encoded into a configurable number of update messages which are then distributed over the course of each second. The consumer receives update messages, and if the messages contain latency information, the consumer decodes them and measures the relative time taken to receive and process the message and its payload.

You can use the Transport API transport-only performance tool (**upacTransportPerf**) to send non-OMM content uni- or bi-directionally. Additionally, this application supports a “reflection” type mode used for round trip measurement.

upacTransportPerf measures latency in all of these configurations, and records independent statistics for each instance of the application.




Figure 4. Transport API C Transport Perf

2. Without a microsecond-resolution synchronization of clocks across machines, the one-way measurement implies that the provider and consumer applications run on the same machine.

2.3 Package Contents

Performance examples are distributed as buildable source code with the Transport API package.




Figure 5. Directory Structure of the Performance Tools

Each example is distributed in its own directory. For Linux and Solaris, **makefiles** are distributed along with the source; on Windows, **vcproj** and solution files are distributed for Visual Studio.

Libxml2, an open source XML-parsing library, is also distributed along with the tools. Each example project is configured to build **Libxml2** as a dependent library.

The **PerfTools/common** directory includes two **.xml** files:

- **350k.xml**: The list of 350,000 items loaded by the consumer (of content published by the non-interactive provider).
- **MsgFile.xml**: The default set of OMM messages.

For more information about examples and their operations, readers can refer to the appropriate application sections in this document. Readers can also refer to the **readme** files and comments included in source.

2.3.1 Building

On Linux and Solaris, use **make** or **gmake** with the included **makefile** to build the examples. To build a 32-bit version, change the **COMPILE_BITS** variable to **32**.

On Windows, open the appropriate project in Visual Studio. Within the package, 64-bit versions of the examples are prefixed with “**_x64**” in the project or solution filename.

2.3.2 Running

On Linux and Solaris, once built, change to the target directory. The **makefile** will have linked the necessary support files (**350k.xml**, **MsgData.xml**, **RDMFieldDictionary**, and **enumtype.def**) to the target directory.

On Windows, after the build process finishes, to run the tool, copy the above support files to the target directory.

2.4 What Is Measured and Reported

2.4.1 Latency

Each performance tool embeds timestamp information in its messages' payloads. The tool uses these timestamps to determine the overall time taken to send and process a message and its payload through the API and, where applicable, the TREP. To ensure that the measurement captures end-to-end latency through the system, the timestamp is taken from the start of the sender's message and payload encoding, and is compared to the time at which the receiver completes its decoding of the message and payload.

When measuring performance, it is important to consider whether or not a particular component acts as a bottleneck on the system. Transport API applications and TREP components provide higher throughput and lower latency than RFA-based applications. In general, Thomson Reuters recommends that you use a Transport API C performance tool to drive and calculate the performance of other non-Transport API C-based performance tools. For example, if you want to test the performance of the RFA C++ consumer, use the Transport API C interactive or non-interactive provider to drive the publishing rather than an RFA C++ providing application.

2.4.2 Throughput and Payload

These tools allow you to control the rate at which messages are sent as well as the content in each message. This allows you to measure throughput and latency using various rates and content, tailored to your specific needs.

2.4.3 Image Retrieval Time

The **upacConsPerf** tool measures the overall time taken to receive a full set of images for items requested through the system. This time is measured from the start of the first request to the reception of the final expected image.

2.4.4 CPU & Memory Usage

Performance tools record a periodic sampling of CPU and Memory usage. This allows for consistent monitoring of resource use and can be used to determine the impact of various features and application modifications.

The CPU and Memory Usage calculations represent the process as a whole, and are not normalized to the number of CPU cores nor running threads. All threads in each tool contribute to the overall execution time; it is possible for the reported CPU usage to be greater than 100% if the application runs multiple busy threads.

2.4.4.1 CPU Usage Calculation

CPU Usage is calculated by periodically querying the OS for applications' "busy" and "overall" times. The difference from the previous value is used to calculate an average CPU usage for each interval and presented as a percentage:

$$\frac{\text{User-space Execution Time} + \text{System-space Execution Time}}{\text{Running Time}}$$

2.4.4.2 Memory Usage Calculation

Memory Usage is calculated using the "Resident Set Size," as provided by the respective operating system (OS). This measures the memory (in RAM) in use by the application.

2.5 Recorded Results and Output

The tools record their test results in the following files:

- Summary File
- Statistics File
- Latency File

2.5.1 Summary File

Each tool records the run's summary to a single file, including:

- The run's configuration
- Overall run results

If you use multiple threads, the file includes results for each thread as well as across all threads. For configuration details, refer to the chapter specific to the application that you use.

An example of recorded summary content for **upacConsPerf** includes the average latency, update rate, and CPU/memory usage for the application's run time.

This summary information is output both to a file and to the console.

2.5.2 Statistics File

Each tool periodically records statistics relevant to that tool. For example, **upacConsPerf** records:

- Latency statistics for updates (and, when so configured, posted content)
- Number of request messages sent and refresh messages received
- Number of update messages received
- Number of generic messages sent and received
- Latency statistics for generic messages (when so configured)

Each tool records these statistics on a per-thread basis. If the tool is configured to use multiple threads, the tool generates a file for each thread. For configuration details, refer to the chapter specific to the application that you use.

Each tool can configure statistics recording via the following options:

- **writeIntervalStats**: The interval (from 1 to n , in seconds) at which timed statistics are written to files and the console.
- **noDisplayStats**: Prevents writing periodic stats to console.

2.5.3 Latency File

You can configure **upacConsPerf** and **upacTransportPerf** to record each individual latency measurement to a file. This is useful for creating plot or distribution graphs, ensuring that recorded latencies are consistent, and for troubleshooting purposes.

These latencies are recorded on a per-thread basis. If the tool is configured to use multiple threads, a file is generated for each thread.

For further details on configuring this behavior, refer to the chapter specific to the application that you use.

Chapter 3 Latency Measurement Details

3.1 Time-slicing

All applications follow a similar model for controlling time: time is divided into small intervals, referred to as “ticks.” During a run, each application has a main loop that runs an iteration once per tick. In this loop, the application performs some periodic action, and then waits until the next tick before starting the loop again.

For example, an application might observe the following loop:

1. Send out a burst of messages.
2. Wait until the time of the next tick. If network notification indicates that any connections have messages available, read them and continue waiting.

Applications can configure this rate using their respective `-tickRate` option. This determines how many ticks occur per second. For example, if you set the tick rate to 100, ticks occur at 10-millisecond intervals.

Applications adjust the message rate to fit the tick rate. For example, if an application wants to send 100,000 messages per second with a tick rate of 100 ticks per second, the application will send 1,000 messages per tick. Adjusting the tick rate affects the smoothness of message traffic by defining the amount of time between bursts:




Figure 6. Time Slicing Algorithm

Depending on the tool, spare time in the tick might be used to perform other actions. For example, after `upacProvPerf` or `upacNIPProvPerf` sends an update burst, the remaining time is used to send outstanding refreshes:




Figure 7. Refresh Publishing Algorithm

Applications always set tick times at fixed intervals as they progress, regardless of what the application does during the interval. For example, if the tick rate is 100 (i.e., 10 ms intervals), and the time of the previous tick was 40ms, then the times of the next ticks are 50 ms, 60 ms, etc... This helps maintain constant overall messaging rates: any irregularities in the timing of the current tick are corrected in subsequent ticks.

3.2 Latency

Latency is measured using timestamps embedded in the messages sent by each application. The receiving application compares this timestamp against the current time to determine the latency.

Each tool sends messages in bursts. To send timestamps, a message is randomly chosen from the message burst and the timestamp is embedded. When this message is received, the receiving application compares it to the current time to determine the latency.




Figure 8. Latency RICs within a Tick

Timestamps are high-resolution and non-decreasing. Because the source of this time varies across platforms and might not be synchronized between multiple machines, update and generic message latency measurements require that the provider and consumer run on the same machine. Posting latency measurements do not require this, as **upacConsPerf** generates both sending and receiving timestamps.

Note: OMM performance tool timestamp information contains the number of microseconds since an epoch.^a **upacTransportPerf** timestamps are provided with nanosecond granularity.

a. Windows uses `QueryPerformanceCounter()`, Solaris uses `gethrtime()`, Linux uses `clock_gettime()` with the monotonic clock.

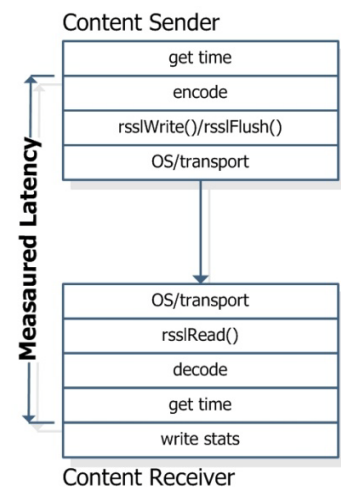


Figure 9. Timing Diagram for Latency Measurements

The standard latency measurement is initiated by the provider, which encodes a starting time into an update. This timestamp is included as a piece of data in the payload using a pre-determined latency FID. On the consumer side, the application processes incoming updates and generic messages, decodes the payload, and looks for updates or generic messages which include the latency FID (known as latency updates). After decoding a latency update or generic message, the consumer takes a second timestamp and compares the two, outputting the difference as the measured latency for that particular update or generic message.

Chapter 4 upacConsPerf

4.1 Overview

A typical OMM consumer application requests content and processes responses to those requests. Thus, the performance consumer makes a large, configurable number of item requests and then processes refresh and update content corresponding to those requests. While processing, the performance consumer decodes all content and collects statistics regarding the count and latency of received messages.

The **upacConsPerf** implements an OMM Consumer using the Transport API C Edition. It connects to a provider (such as **upacProvPerf** or TREP), requests items, and processes the refresh and update messages it receives, calculating statistics such as update rate and latency. Additionally, the consumer can send post messages through the system at a configured rate, measuring the round-trip latency of posted content.

At startup, the consumer performs some administrative tasks, such as logging into the system, obtaining a source directory, and maybe requesting a dictionary. After the consumer is satisfied that the correct service is available and that the provider is accepting requests, the consumer begins requesting data. **upacConsPerf** uses Transport API Value-Add Administration Domain Representations to complete its start-up tasks. For more information, refer to the *Transport API Value Added Components Developers Guide*.

4.2 Threading and Scaling

The Transport API is designed to allow calls from multiple threads, such that applications can scale their work across multiple cores. Applications can leverage this feature by creating multiple threads to handle multiple connections through the Transport API. As such, each application enables global locking when calling `rsslInitialize`.

Configure **upacConsPerf** for multiple threads using the `-threads` command-line option. When multiple threads are configured, each thread opens its own connection to the provider. **upacConsPerf** divides its list of items among the threads (you can use the command line option, `-commonItemCount`, to request the same type and number of items on all connections).

The main thread monitors the other threads and collects and reports statistics from them.

4.2.1 Consumer Lifecycle

The lifecycle of **upacConsPerf** is divided into the following sections:




Figure 10. upacConsPerf Lifecycle

1. Application and Transport API Initialization.

upacConsPerf loads its configuration, initializes the Transport API, loads its item list using the specified file, and starts the thread(s) which connect to the provider to perform the test.

- The main thread periodically collects and writes statistics from the connection thread(s) until the test is over. All subsequent steps are performed by each thread.
- Connection: the connection thread connects to the provider. If the connection fails, it continually attempts to reconnect until the connection succeeds. When the connection succeeds, the test begins and any subsequent disconnection ends the test.
- Login: the connection thread provides its login requests and waits for the provider's response.
- Directory: the connection thread opens a directory stream and searches for the configured service name.
- Startup state: when the service is available, the "startup" phase of the performance measurement begins. During this phase, the connection thread continually performs the following actions:
 - Sends bursts of requests, until all desired items have been requested.
 - Reads from the transport, processing refresh, update, and generic message traffic from the provider.

The "startup" phase continues until all items receive a refresh containing an Open/OK state. All latency statistics recorded up to this point are reported as "startup" statistics.

2. Steady state.

The connection thread continually performs the following actions:

- If configured for posting, the thread sends a burst of post messages.
- Reads from the transport, processing updates from the provider.
- If configured to do so, sends a burst of generic messages.

The "steady state" phase continues for the period of time specified in the command line. Latency statistics recorded during this phase are reported as "steady state" statistics.

3. Application shutdown and cleanup.

The connection thread disconnects and stops. The main thread collects all remaining information from the connection threads, cleans them up, and writes the final summary statistics. The main thread then uninitializes the Transport API, any remaining resources, and exits.

4.2.2 Diagram

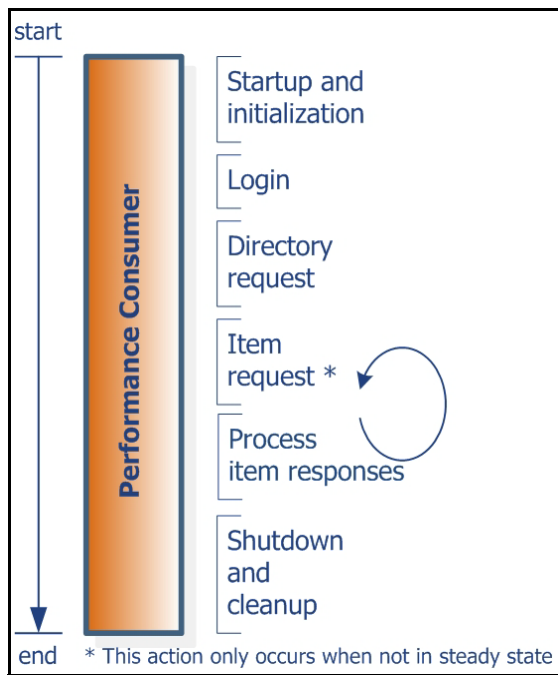


Figure 11. upacConsPerf Application Flow

4.3 Latency Measurement

Provider applications encode the timestamp as part of their message payload. The initial timestamp is taken at the start of encoding, and added as field TIM_TRK_1 (**3902**). When this field is detected, the **upacConsPerf** gets the current time and computes the difference to measure latency.

4.3.1 Consumer Latency

► Consumer Latency Measurement Sequence:

1. Read the message from the API (received via the underlying transport).
2. Decode the message.
3. Check whether the payload contains latency information, if so:
 - Get the current time (**t2**).
 - Calculate the difference between timestamps.
 - Store the result as part of the recorded output information.

4.3.2 Posting Latency

You can configure **upacConsPerf** to send on-stream posts in which case the consumer periodically sends bursts of post messages for specified items in the item list file. You can also configure the tool to include latency information in its posts. When configured in this manner, **upacConsPerf** adds latency information to random post messages. When the posted content returns on the stream, **upacConsPerf** decodes the timestamp and measures the difference to determine posting latency.

► Posting Latency Measurement Sequence:

1. Get the current time (**t1**).
2. Obtain an output buffer using `rsslGetBuffer()`.
3. Encode the message, including the time (**t1**).
4. Pass the message to the API, which then passes it to the underlying transport.
5. When processing received content, check to see whether the payload contains latency information, if so:
 - Get the current time (**t2**).
 - Calculate the difference between timestamps.
 - Store the result in the recorded output information.

The time at the start of encoding is encoded as a timestamp in the payload as field TIM_TRK_2 (**3903**). When the payload from the post returns from the platform, the consumer compares the timestamp to the current time to determine the posting latency.

4.4 upacConsPerf Configuration Options

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-commonItemCount	0	If multiple consumer threads are created (see -threads), each thread normally requests a unique set of items on its connection. This option specifies the number of common items to be requested by all connections.
-connType	socket	Specifies the consumer's connection type. upacConsPerf supports the following options: socket , http , and encrypted .
-genericMsgLatencyRate	0	Controls the number of generic messages sent that contain latency information. This must be greater than the tick rate (see -tickRate) and less than the total generic message rate (see -genericMsgRate).
-genericMsgRate	0	Controls the number of generic messages sent per second. This cannot be less than the tick rate, unless it is zero.
-h	localhost	Specifies the hostname to which the consumer connects.
-if	<none>	Configures interfaceName (an RsslConnectOptions parameter), which configures the network interface card (NIC) through which the consumer makes its connection. If your machine straddles networks, you can use this setting to force the consumer to use a particular network.
-inputBufs	15	Configures numInputBufs (an RsslConnectOptions parameter) which configures the size of the Transport API's input queue. Use a setting large enough to accommodate incoming data.
-itemCount	100000	Sets the total number of items requested by the consumer.
-itemFile	350k.xml	Configures the name of the item list file.
-latencyFile	<none>	Specifies the name of the log file in which upacConsPerf logs the latency retrieved from individual latency updates, generic messages, and posts. If a name is not specified, logging is disabled.
-msgFile	MsgData.xml	Configures the name of the file used by the consumer to determine the makeup of message payloads when posting (see -postingRate).
-noDisplayStats	(no argument)	Turns off printing statistics to the screen.
-outputBufs	5000	Configures guaranteedOutputBuffers (an RsslConnectOptions parameter) which specifies the minimum guaranteed number of output buffers created for each channel. Setting this parameter to an appropriate size can aid performance if the consumer is posting messages (see -postingRate). You should configure enough buffers so that the provider does not run out of buffers while writing, but at the same time limit the number so as to conserve memory and optimize performance.
-p	14002	Specifies the port number to which the consumer connects.
-postingLatencyRate	0	Controls the number of posts sent per second that contain latency information. This must be greater than the tick rate (see -tickRate) and less than the total post message rate (see -postingRate).

Table 2: upacConsPerf Configuration Options

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-postingRate	0	Configures the consumer for posting. Sets the number of posting messages the consumer sends, per second. This cannot be less than the tick rate, unless it is zero.
-reactor	(no argument)	Use the Value Added Reactor instead of the Transport API channel for sending and receiving. For details on Value Added Components, refer to the <i>Transport API Value Added Components Developers Guide</i> .
-recvBufSize	<none>	Sets the size (in bytes) of the system receive buffer. When unspecified, the OS setting is used.
-requestRate	500000	Sets the number of item requests sent (per second).
-sendBufSize	<none>	Sets the size (in bytes) of the system send buffer. When unspecified, the OS setting is used.
-serviceName	DIRECT_FEED	Configures the name of the service used by the consumer to request items. The consumer begins requesting items whenever this service is found and appears ready.
-snapshot	(no argument)	Opens all items as snapshots, even if not specified in the item list file, and exits upon receiving all the solicited images. This is different from setting -steadyStateTime to 0 in that the requests are specifically made without the "RSSL_RQF_STREAMING" RequestMsg flag.
-statsFile	ConsStats	Configures the base name that the consumer uses when writing its test statistics.
-steadyStateTime	300	Configures how long (in seconds) the consumer continues to run the test after receiving the last expected image. steadyStateTime has a second function: after beginning the test, if the consumer does not receive all expected images within this segment of time, the consumer times out. In this case, it exits and indicates that it did not reach steady state.
-summaryFile	ConsSummary.out	Configures the name of the file to which the consumer writes its test summary.
-threads	<none>	Sets the number of threads the consumer starts. Each specified thread starts its own connection to the configured provider.
-tickRate	1000	Sets the number of 'ticks' per second (the number of times per second the main loop of the consumer occurs). Adjusting the tick rate changes the size of request/post bursts; a higher tick rate results in smaller individual bursts, creating smoother traffic.
-watchlist	(no argument)	Use the Value Added Reactor watchlist instead of the Transport API channel for sending and receiving. For details on Value Added Components, refer to the <i>Transport API Value Added Components Developers Guide</i> .
-writeStatsInterval	5	Configures the frequency (in seconds) at which statistics are printed to the screen and statistics file.

Table 2: upacConsPerf Configuration Options (Continued)

4.5 Input

upacConsPerf requires the following files:

- Dictionary files to validate fields in the message data. **RDMFieldDictionary** and **enumtype.def** are provided with the package.
- An XML file that describes the items that **upacConsPerf** should request and (when configured) which items to post. The package includes a default file (**350k.xml**).
- If the consumer is posting, an XML file that describes post message data. The package includes a default file (**MsgData.xml**) with this information.

For more details on input file information, refer to Chapter 9.

4.6 Output

upacConsPerf records statistics during a test such as:

- Item requests sent and images received
- Image retrieval time
- The update rate
- The post message rate
- The generic message rate
- Latency statistics
- CPU and memory usage

For more details on output file information, refer to Chapter 10.

4.6.1 upacConsPerf Summary File Sample

```

--- TEST INPUTS ---

Steady State Time: 90
  Connection Type: socket
    Hostname: 192.168.12.109
    Port: 14002
    Service: DIRECT_FEED
  Thread List: 7
  Output Buffers: 2048
  Input Buffers: 2048
  Send Buffer Size: 0(use default)
  Recv Buffer Size: 0(use default)
  High Water Mark: 0(use default)
  Interface Name: (use default)
  Tcp_NoDelay: Yes
  Username: (use system login name)
  Item Count: 100000
Common Item Count: 0
  Request Rate: 500000
Request Snapshots: No
  Posting Rate: 0
Latency Posting Rate: 0
  Item File: 350k.xml
  Data File: MsgData.xml
  Summary File: ConsSummary.out
  Stats File: ConsStats
  Latency Log File: latencyLog
  Tick Rate: 1000

--- OVERALL SUMMARY ---

Startup State Statistics:
  Sampling duration (sec): 3.826
  Latency avg (usec): 663.2
  Latency std dev (usec): 693.1
  Latency max (usec): 2242.0
  Latency min (usec): 138.0
  Avg update rate: 100009

Steady State Statistics:
  Sampling duration (sec): 90.800
  Latency avg (usec): 157.5
  Latency std dev (usec): 31.8
  Latency max (usec): 614.0
  Latency min (usec): 106.0
  Avg update rate: 100000

```

Overall Statistics:

```

Sampling duration (sec): 94.626
Latency avg (usec): 177.8
Latency std dev (usec): 172.0
Latency max (usec): 2242.0
Latency min (usec): 106.0
CPU/Memory samples: 19
CPU Usage max (%): 21.02
CPU Usage min (%): 0.00
CPU Usage avg (%): 1.11
Memory Usage max (MB): 62.57
Memory Usage min (MB): 56.65
Memory Usage avg (MB): 57.87

```

Test Statistics:

```

Requests sent: 100000
Refreshes received: 100000
Updates received: 9462101
Image retrieval time (sec): 3.826
Avg image rate: 26139
Avg update rate: 100000

```

Code Example 1: upacConsPerf Summary File Sample**4.6.2 upacConsPerf Statistics File Sample**

```

UTC, Latency updates, Latency avg (usec), Latency std dev (usec), Latency max (usec), Latency min (usec),
    Images, Update rate(msg/sec), Posting Latency updates, Posting Latency avg (usec), Posting
    Latency std dev (usec), Posting Latency max (usec), Posting Latency min (usec), CPU usage (%),
    Memory(MB)
2012-11-13 01:44:48, 49, 550.1, 644.5, 2242.0, 138.0, 100000, 92340, 0, 0.0, 0.0, 0.0, 0.0, 21.02, 56.65
2012-11-13 01:44:53, 48, 157.2, 14.7, 193.0, 136.0, 0, 100018, 0, 0.0, 0.0, 0.0, 0.0, 0.04, 56.69
2012-11-13 01:44:58, 51, 173.2, 84.0, 614.0, 106.0, 0, 100001, 0, 0.0, 0.0, 0.0, 0.0, 0.00, 56.69
2012-11-13 01:45:03, 49, 155.0, 18.6, 190.0, 110.0, 0, 100000, 0, 0.0, 0.0, 0.0, 0.0, 0.00, 56.69
2012-11-13 01:45:08, 52, 154.3, 18.9, 188.0, 110.0, 0, 100000, 0, 0.0, 0.0, 0.0, 0.0, 0.02, 56.69
2012-11-13 01:45:13, 48, 158.2, 15.5, 189.0, 111.0, 0, 100002, 0, 0.0, 0.0, 0.0, 0.0, 0.00, 56.69
2012-11-13 01:45:18, 51, 167.0, 62.9, 482.0, 114.0, 0, 100017, 0, 0.0, 0.0, 0.0, 0.0, 0.00, 56.69

```

Code Example 2: upacConsPerf Statistics File Sample

4.6.3 upacConsPerf Latency File Sample

```

Message type, Send Time, Receive Time, Latency(usec)
Upd, 2081854121264, 2081854121654, 390
Upd, 2081854192351, 2081854194582, 2231
Upd, 2081854246268, 2081854248510, 2242
Upd, 2081854301286, 2081854302685, 1399
Upd, 2081854336328, 2081854338061, 1733
Upd, 2081854387298, 2081854387494, 196
Upd, 2081854495308, 2081854495563, 255
Upd, 2081854608277, 2081854608456, 179
Upd, 2081854642285, 2081854642437, 152
Upd, 2081854859315, 2081854859482, 167
Upd, 2081855120275, 2081855120423, 148
Upd, 2081855123318, 2081855124492, 1174
Upd, 2081855125345, 2081855125558, 213
Upd, 2081855292310, 2081855292452, 142
Upd, 2081855834263, 2081855835084, 821
Upd, 2081855946300, 2081855946479, 179

```

Code Example 3: upacConsPerf Latency File Sample

4.6.4 upacConsPerf Console Output Sample

```

005: Images: 100000, UpdRate: 92340, CPU: 21.02%, Mem: 56.65MB
    Latency (usec): Avg: 550.1 StdDev: 644.5 Max: 2242 Min: 138, Msgs: 49
    - Image retrieval time for 100000 images: 3.826s (26139 images/s)
010: Images: 0, UpdRate: 100018, CPU: 0.04%, Mem: 56.69MB
    Latency (usec): Avg: 157.2 StdDev: 14.7 Max: 193 Min: 136, Msgs: 48
015: Images: 0, UpdRate: 100001, CPU: 0.00%, Mem: 56.69MB
    Latency (usec): Avg: 173.2 StdDev: 84.0 Max: 614 Min: 106, Msgs: 51
020: Images: 0, UpdRate: 100000, CPU: 0.00%, Mem: 56.69MB
    Latency (usec): Avg: 155.0 StdDev: 18.6 Max: 190 Min: 110, Msgs: 49
025: Images: 0, UpdRate: 100000, CPU: 0.02%, Mem: 56.69MB
    Latency (usec): Avg: 154.3 StdDev: 18.9 Max: 188 Min: 110, Msgs: 52
030: Images: 0, UpdRate: 100002, CPU: 0.00%, Mem: 56.69MB
    Latency (usec): Avg: 158.2 StdDev: 15.5 Max: 189 Min: 111, Msgs: 48
035: Images: 0, UpdRate: 100017, CPU: 0.00%, Mem: 56.69MB
    Latency (usec): Avg: 167.0 StdDev: 62.9 Max: 482 Min: 114, Msgs: 51

```

Code Example 4: upacConsPerf Console Output Sample

Chapter 5 upacProvPerf

5.1 Overview

A typical interactive provider allows consuming applications, including TREP, to connect. Once connected, consumers log in and request content. The interactive provider will respond, providing requested content when possible and a status indicating some type of failure when not possible. While a provider in a production environment might get its data from an external source or by performing a calculation on some other data, the performance provider generates its data internally.

upacProvPerf implements an OMM Interactive Provider using the Transport API. It starts a server which allows OMM consumers to connect (either directly or through TREP), and provides customizable refresh messages and update messages for requested items.

When a new connection is being established, the provider performs some administrative tasks, such as processing login messages, handling directory requests, and (optionally) providing a dictionary. This application uses the Transport API Value Add Administration Domain Representations to complete these tasks. For more information, refer to the *Transport API Value Added Components Developers Guide*.

5.2 Threading and Scaling

The Transport API is designed to allow calls from multiple threads, such that applications can scale their work across multiple cores by creating multiple threads to handle multiple connections through the Transport API. To support this multi-threading feature, each application must enable global locking when calling `rsslInitialize`.

You can configure **upacProvPerf** for multiple threads by using the `-threads` command-line option. When multiple threads are configured, consumer connections are balanced such that each thread receives an equal number of connections.

The main thread monitors the other threads and collects and reports statistics from them.

5.3 Provider Lifecycle

The lifecycle of **upacProvPerf** is divided into the following sections:




Figure 12. upacProvPerf Lifecycle

1. Application and Transport API Initialization.

upacProvPerf loads its configuration, initializes the Transport API, loads its dictionary and sample message data using specified files, and starts one or more threads (as configured) to provide data to consumers.

The main thread has two roles:

- Accept consumer connections and pass them to one of the provider threads.
- Periodically collect and write statistics from the connection thread(s) until the test is over.

2. Handle Login, Directory, and Item Requests; Provide Updates.

The provider thread performs the following actions continually until its run time expires:

- Add new connections passed from the main thread.
- Send a burst of updates for items currently open on existing connections.
- Send a burst of generic messages (if configured to do so).
- Use available spare time to provide images for items that need them.
- Use available spare time to read from the transport, processing any Login, Directory, or Item requests.

3. Shutdown and cleanup.

The provider thread stops. The main thread collects any remaining data from the connection threads, cleans them up, and writes the final summary statistics. The main thread then cleans up the Transport API and remaining resources, and exits.

upacProvPerf should run long enough to allow connected consumers to complete their measurements.




Figure 13. upacProvPerf Application Flow

5.4 Latency Measurement

upacProvPerf encodes the timestamp as part of its message payload. The timestamp is taken at the start of encoding and added as field TIM_TRK_1 (3902). Latency is measured after **upacConsPerf** completes decoding.

► Interactive Provider Latency Measurement Sequence:

1. Get the current time (**t1**).
2. Obtain an output buffer using **rsslGetBuffer**.
3. Encode the message, including time **t1**.
4. Pass the message to the API, which passes it to the underlying transport.
5. The consuming application receives the timestamp in the payload and compares it against the current time to calculate latency.

5.5 upacProvPerf Configuration Options

upacProvPerf uses the following command line options:

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-directWrite	(no argument)	Sets whether to use the RSSL_WRITE_DIRECT_SOCKET_WRITE flag when calling rsslWrite . This flag causes the write to attempt to bypass the Transport API output queue, reducing latency but at a cost of throughput capability.
-genericMsgLatencyRate	0	Sets the number of generic messages sent (per second) that contain latency data. This number must be greater than the tick rate (see -tickRate) and less than the total generic message rate (see -genericMsgRate).
-genericMsgRate	0	Sets the number of generic messages sent per second. This number cannot be less than the tick rate (unless it is zero).
-highWaterMark	6000	Configures the quantity of data (in bytes) that the Transport API queues before automatically flushing it to the network. Adjusting this might provide a tradeoff of throughput vs. latency when writing large bursts of data.
-if	<none>	Sets interfaceName (an RsslBindOptions parameter), which configures the NIC that the provider uses for its server. If your machine straddles networks, you can use this setting to force the provider to use a particular network.
-latencyFile	<none>	Specifies the name of the log file in which upacProvPerf logs the latency retrieved from individual latency updates, generic messages, and posts. If a name is not specified, logging is disabled.
-latencyUpdateRate	10	<p>Sets the number of updates sent per second containing latency information.</p> <p>Note: You must use a setting greater than the tick rate (see -tickRate) and less than the total update rate (see -updateRate)</p>

Table 3: upacProvPerf Configuration Options

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-maxFragmentSize	6144	Sets maxFragmentSize (an RsslBindOptions parameter) which specifies the size of RsslBuffers. Messages less than this size are sent as a single buffer while larger messages will be fragmented and reassembled into multiple buffers by the Transport API.
-maxPackCount	1	Sets the number of buffers packed when sending refresh and update bursts. A setting greater than 1 causes the provider to use the rsslPackBuffer function to pack messages, raising the maximum throughput capability (at a potential cost to latency). When packing, the provider attempts to pack the number of buffers specified by -maxPackCount into a buffer of the length specified by -packBufSize . The provider periodically displays averages of how many messages are successfully packed.
-msgFile	MsgData.xml	Specifies the file that the provider uses to determine message content.
-noDisplayStats	(no argument)	Turns off printing statistics to the screen.
-openLimit	1000000	Configures the maximum number of items the provider allows each client to open.
-outputBufs	5000	Sets guaranteedOutputBuffers (an RsslBindOptions parameter) which specifies the minimum guaranteed number of output buffers created for each channel that a server accepts. Set this large enough so that the provider does not run out of buffers while writing, but not so large that it wastes memory and slows performance.
-p	14002	Specifies the port on which upacProvPerf listens for connections.
-packBufSize	6000	Configures the size (in bytes) of the buffer used for packing (used in conjunction with -maxPackCount). For more information, refer to Section 11.2.3 and the <i>Transport API Developers Guide</i> .
-reactor		Send and receive using the Value Added Reactor instead of the Transport API channel. For details on using Value Added Components, refer to the <i>Transport C Edition API Value Added Components Developers Guide</i> .
-recvBufSize	<none>	Configures the size (in bytes) of the system's receive buffer size. When unspecified, the OS setting is used.
-refreshBurstSize	10	After the provider completes an update burst, it uses the time before the next burst to send any needed refreshes, monitoring the time to see whether the next tick time has been reached. This option configures how often the provider checks the time (in case checking is expensive for the system).
-runTime	360	Sets the length of time upacProvPerf runs (in seconds).
-sendBufSize	<none>	Configures the size of the system's send buffer. When unspecified, the OS setting is used.
-serviceld	1	Specifies the ID of the provider's service.
-serviceName	DIRECT_FEED	Specifies the name of the provider's service.

Table 3: upacProvPerf Configuration Options (Continued)

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-statsFile	ProvStats	Specifies the base name used to write the provider's test statistics.
-summaryFile	ProvSummary.out	Specifies the base file name used to write the provider's test summary.
-tcpDelay	(no argument)	Configures tcp_nodelay (an RsslBindOptions parameter) which sets whether the underlying connection uses Nagle's algorithm , a method for batching data and optimizing network bandwidth. By default, performance tools disable Nagle's algorithm due to the increased latency from the buffering, but this option enables it, which can raise throughput capability.
-threads	<none>	Sets the number of threads that the provider starts and the CPU core to which each thread binds. Each thread acts as a provider, and connected consumers are balanced among all threads. Example: "1,3" creates two threads to handle consumer connections, respectively bound to CPU cores 1 and 3.
-tickRate	1000	Sets the number of 'ticks' (cycles completed by the provider's main loop) per second. Adjusting the tick rate changes the size of update bursts: higher tick rates result in smaller individual bursts, creating smoother traffic.
-updateRate	100000	Configures the number of updates sent per second, per connection. Note: This cannot be less than the tick rate, unless it is zero.
-writeStatsInterval	5	Sets how often statistics are printed to the screen and statistics file (in seconds).

Table 3: upacProvPerf Configuration Options (Continued)

5.6 Input Files

upacConsPerf requires the following files:

- Dictionary files to validate fields in the message data. **RDMFieldDictionary** and **enumtype.def** are provided with the package.
- An XML file that describes the items that **upacConsPerf** should request and (when configured) which items to post. The package includes a default file (**350k.xml**).
- If the consumer is posting, an XML file that describes post message data. The package includes a default file (**MsgData.xml**) with this information.

For more details on input file information, refer to Chapter 9.

5.7 Output

upacProvPerf records statistics during a test such as:

- Item requests received
- Updates sent
- Posts received and reflected
- CPU and memory usage

For more detailed output file information, refer to Chapter 10.

5.7.1 upacProvPerf Summary File Sample

```

--- TEST INPUTS ---

        Run Time: 180
        Port: 18003
    Thread List: 5
    Output Buffers: 2048
Max Fragment Size: 6144
    Send Buffer Size: 0(use default)
    Recv Buffer Size: 0(use default)
    Interface Name: 192.168.10.114
    Tcp_NoDelay: Yes
    Tick Rate: 1000
    Use Direct Writes: No
    High Water Mark: 1000 bytes
    Summary File: ProvSummary.out
    Stats File: ProvStats
Write Stats Interval: 5
    Display Stats: Yes
    Update Rate: 100000
Latency Update Rate: 10
    Refresh Burst Size: 10
        Data File: MsgData.xml
        Packing: No
        Service ID: 1
    Service Name: DIRECT_FEED
    OpenLimit: 1000000

--- OVERALL SUMMARY ---

Overall Statistics:
    Image requests received: 100000
    Posts received: 0
    Updates sent: 16785700
    Posts reflected: 0
    CPU/Memory samples: 37
    CPU Usage max (%): 5.56

```

```

CPU Usage min (%): 0.00
CPU Usage avg (%): 0.21
Memory Usage max (MB): 42.29
Memory Usage min (MB): 13.34
Memory Usage avg (MB): 39.51

```

Code Example 5: upacProvPerf Summary File Sample

5.7.2 upacProvPerf Statistics File Sample

```

UTC, Requests received, Images sent, Updates sent, Posts reflected, CPU usage (%), Memory(MB)
2012-11-13 01:44:35, 0, 0, 0, 0, 0.08, 13.34
2012-11-13 01:44:40, 0, 0, 0, 0, 0.00, 13.39
2012-11-13 01:44:45, 49664, 49664, 185670, 0, 2.20, 24.76
2012-11-13 01:44:50, 50336, 50336, 500013, 0, 5.56, 36.28
2012-11-13 01:44:55, 0, 0, 499995, 0, 0.00, 36.28
2012-11-13 01:45:00, 0, 0, 500003, 0, 0.00, 36.28
2012-11-13 01:45:05, 0, 0, 499990, 0, 0.00, 40.23

```

Code Example 6: upacProvPerf Statistics File Sample

5.7.3 upacProvPerf Console Output Sample

```

005: UpdRate:      0, CPU:    0.08%, Mem:   13.34MB
010: UpdRate:      0, CPU:    0.00%, Mem:   13.39MB
015: UpdRate:    37134, CPU:    2.20%, Mem:   24.76MB
    - Received 49664 item requests (total: 49664), sent 49664 images (total: 49664)
020: UpdRate:   100002, CPU:    5.56%, Mem:   36.28MB
    - Received 50336 item requests (total: 100000), sent 50336 images (total: 100000)
025: UpdRate:    99999, CPU:    0.00%, Mem:   36.28MB
030: UpdRate:   100000, CPU:    0.00%, Mem:   36.28MB
035: UpdRate:    99998, CPU:    0.00%, Mem:   40.23MB

```

Code Example 7: upacProvPerf Console Output Sample

Chapter 6 upacNIPProvPerf

6.1 Overview

A Non-Interactive Provider (NIP) publishes content regardless of consumer requests by connecting to an ADH and publishing content to the ADH cache. After login, an NIP publishes a service directory and then starts sending data for supported items.

upacNIPProvPerf implements an OMM NIP using the Transport API C Edition for use with the Advanced Data Hub (ADH) on the TREP. It connects and logs into an ADH, publishes its service, and then provides images and updates.

When connecting, the NIP performs some administrative tasks, like processing system logins and publishing a directory refresh. The **upacNIPProvPerf** uses Transport API Value Add Administration Domain Representations to complete these tasks. For more information, refer to the *Transport API Value Added Components Developers Guide*.

6.2 Threading and Scaling

You can configure **upacNIPProvPerf** for multiple threads via the **-threads** command-line option. When you configure multiple threads, each thread opens its own connection to the ADH, and the list of items is divided among all threads. You can use the **-commonItemCount** option to control the number of items that will be sent across all threads.

The main thread monitors the other threads and then collects and reports their statistics.

6.3 Non-Interactive Provider Lifecycle




Figure 14. upacNIPProvPerf Lifecycle

The lifecycle of **upacNIPProvPerf** is divided into the following sections:

1. Application and Transport API Initialization.

In this phase **upacNIPProvPerf**:

- Loads its configuration.
- Initializes the Transport API.
- Loads its dictionary, item list, and sample message data using the specified files.

- Starts the thread(s) that will connect to the ADH to perform the test.
 - The main thread begins cycling: periodically collecting and writing statistics from the connection thread(s).
 - Connection threads connect to the ADH. If a connection fails, the thread continually attempts to reconnect until the connection succeeds. Once the connection succeeds, the test begins and any subsequent disconnection ends the test.

2. ADH Login and Directory.

The connection thread sends login requests and waits for the ADH response. After a successful login, the connection thread publishes its service, followed by an item image to begin the publishing phase.

3. Provide OMM content.

The connection thread begins providing the items specified in its item list, continually performing the following actions:

- Send a burst of updates for open items.
- If refreshes are needed, use spare time in the tick to send them.
- Using any spare time left, read from the transport and process incoming messages.

4. Application shutdown and cleanup.

The connection thread disconnects and stops. The main thread collects any remaining information from the connection threads, cleans them up, and writes the final summary statistics. The main thread then cleans up the Transport API and any remaining resources and then exits.

Run **upacNIProvPerf** for a long enough period of time to allow for connected consumers to complete their measurements.




Figure 15. upacNIProvPerf Application Flow

6.4 Latency Measurement

The Transport API is designed to allow calls from multiple threads, such that applications can scale their work across multiple cores. Applications can take advantage of this by creating multiple threads to handle multiple connections through the Transport API. To support this multi-threading, each application enables global locking when calling `rsslInitialize`.

upacNIProvPerf encodes a timestamp as part of its message payload. The timestamp is taken at the start of encoding and added as field TIM_TRK_1 (3902). Latency is measured after **upacConsPerf** decodes the message and payload.

► Non-Interactive Provider Latency Measurement Sequence:

1. Get the current time (**t1**).
2. Obtain an output buffer using `rsslGetBuffer`.
3. Encode the message, including time **t1**.
4. Pass the message to the API, which passes it to underlying transport.
5. The consuming application receives a timestamp in the payload and compares it to the current time to calculate latency.

6.5 upacNIProvPerf Configuration Options

upacProvPerf uses the following command line options:

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-connType	socket	Specifies the connection type used by the application. Supported options are socket , http , encrypted , and reliableMCast .
-directWrite	(no argument)	Sets whether the RSSL_WRITE_DIRECT_SOCKET_WRITE flag is used when calling <code>rsslWrite</code> . Using this flag causes the write to attempt to bypass the Transport API output queue, reducing latency but at a cost of throughput capability.
-h	localhost	When using TCP socket connection types, specifies the hostname of the machine on which the ADH (to which the provider connects) runs.
-highWaterMark	6000	Configures the Transport API's "High-water Mark," the amount of data (in bytes) that the Transport API queues before flushing it to the network. Adjusting this might provides a trade-off of throughput vs. latency when writing large bursts of data.
-if	<none>	Configures interfaceName (an RsslConnectOptions parameter), which configures the NIC that the provider uses for its server. On computers connected to multiple networks, you can use this parameter to force the provider to use the desired network.
-itemFile	350k.xml	Specifies the file that contains a list of items the provider will publish.
-latencyUpdateRate	10	Sets the number of updates with latency information sent per second.
		Note: This must be greater than the tick rate (see <code>-tickRate</code>) but less than the total update rate (see <code>-updateRate</code>).

Table 4: upacNIProvPerf Configuration Options

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-maxPackCount	1	Sets the number of buffers packed when sending refresh and update bursts. Specifying a count greater than 1 causes the provider to use the rsslPackBuffer function to pack messages, raising maximum throughput capability but potentially having a negative affect on latency. When packing, the provider attempts to pack the number of buffers specified by -maxPackCount into a buffer of the length specified by -packBufSize . The provider periodically displays averages of how many messages are successfully packed.
-msgFile	MsgData.xml	Specifies the file that determines the provider's message content.
-noDisplayStats	(no argument)	Turns off printing statistics to the screen.
-outputBufs	5000	Configures guaranteedOutputBuffers (an RsslConnectOptions parameter), which sets the minimum guaranteed number of output buffers created for each channel. Set this value large enough so that the provider does not run out of buffers while writing but low enough so as not to waste memory and slow performance.
-p	14003	When using TCP socket connection types, specifies the port number the provider uses to connect to the ADH.
-packBufSize	6000	Specifies the size of the buffer (in bytes) used for packing. Used in conjunction with -maxPackCount .
-ra	<none>	When using a reliable multicast connection, configures the multicast receive address.
-reactor		Send and receive using the Value Added Reactor instead of the Transport API channel. For details on using Value Added Components, refer to the <i>Transport C Edition API Value Added Components Developers Guide</i> .
-recvBufSize	<none>	Configures the system receive buffer size. By default, the OS setting is used.
-refreshBurstSize	10	After the provider completes an update burst, it uses the time before the next burst to send any needed refreshes, monitoring the time to see whether it is time for the next tick time. This option configures how often the provider checks the time (in case checking is expensive for the system).
-rp	<none>	When using reliable multicast connection type, configures the multicast receive port.
-runTime	360	Sets the length of time for which upacNIProvPerf runs, in seconds.
-sa	<none>	When using reliable multicast connection type, configures the multicast send address.
-sendBufSize	<none>	Configures the system send buffer size. By default, the OS setting is used.
-serviceId	1	Specifies the provider's service ID.

Table 4: upacNIProvPerf Configuration Options (Continued)

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-serviceName	DIRECT_FEED	Specifies the provider's service name.
-sp	<none>	When using reliable multicast connection type, configures the multicast send port.
-statsFile	NIProvStats	Specifies the base filename used to write the provider's test statistics.
-summaryFile	NIProvSummary.out	Specifies the base filename used to write the provider's test summary.
-tcpDelay	(no argument)	Configures tcp_nodelay (an RsslConnectOptions parameter), an option that sets whether the underlying connection uses Nagle's algorithm; a method for more efficiently using network packet headers by batching data. By default, upacNIProvPerf disables Nagle's algorithm due to the increased latency from the buffering, but -tcpDelay enables it, which can raise throughput capability.
-tickRate	1000	Sets the number of ticks per second (the number of cycles per second made by the provider's main loop). Adjusting the tick rate changes the size of update bursts; higher tick rates result in smaller individual bursts and smoother traffic.
-u	<none>	When using reliable multicast connection type, sets the unicast port.
-updateRate	100000	Sets the total number of updates sent per second, per connection. Note: This cannot be less than the tick rate, unless it is 0.
-writeStatsInterval	5	Sets how often upacNIProvPerf prints statistics to the screen and statistics file.

Table 4: upacNIProvPerf Configuration Options (Continued)

6.6 Input Files

upacNIProvPerf requires the following files:

- An XML file that describes **upacNIProvPerf** message data. By default, the package includes the file: **MsgData.xml**.
- Dictionary files to validate fields present in the message data. By default, the package includes the **RDMFieldDictionary** and **enumtype.def** files.
- An XML file that describes the items that **upacNIProvPerf** should publish. By default, the package includes the file, **350k.xml**.

For more detailed input file information, refer to Chapter 9.

6.7 Output

upacNIProvPerf records statistics during a test, such as:

- The number of sent images
- The number of sent updates
- CPU and memory usage

For more detailed output file information, refer to Chapter 10.

6.7.1 upacNIProvPerf Summary File Sample

```

--- TEST INPUTS ---

        Run Time: 60 sec
    Connection Type: socket
        Hostname: 192.168.3.122
            Port: 14003
        Thread List: 5
    Output Buffers: 6000
Max Fragment Size: 6144
    Send Buffer Size: 0(use default)
    Recv Buffer Size: 0(use default)
    High Water Mark: 0(use default)
    Interface Name: 192.168.3.190
        Username: (use system login name)
        Tcp_NoDelay: Yes
        Item Count: 100000
    Common Item Count: 0
        Tick Rate: 1000
    Use Direct Writes: No
        Summary File: NIProvSummary.out
        Stats File: NIProvStats
Write Stats Interval: 5
    Display Stats: Yes
        Update Rate: 100000
Latency Update Rate: 10
    Refresh Burst Size: 10
        Item File: 350k.xml
        Data File: MsgData.xml
        Packing: No
        Service ID: 1
        Service Name: NI_PUB

--- OVERALL SUMMARY ---

Overall Statistics:
    Images sent: 100000
    Updates sent: 6100000
    CPU/Memory samples: 12
    CPU Usage max (%): 12.94
    CPU Usage min (%): 0.00
    CPU Usage avg (%): 1.08
    Memory Usage max (MB): 69.46
    Memory Usage min (MB): 69.44
    Memory Usage avg (MB): 69.46

```

Code Example 8: upacNIProvPerf Summary File Sample**6.7.2 upacNIProvPerf Statistics File Sample**

```

UTC, Images sent, Updates sent, CPU usage (%), Memory (MB)
2012-11-14 19:32:08, 100000, 503706, 12.94, 69.44
2012-11-14 19:32:13, 0, 507003, 0.00, 69.46
2012-11-14 19:32:18, 0, 507016, 0.00, 69.46
2012-11-14 19:32:23, 0, 506980, 0.00, 69.46
2012-11-14 19:32:28, 0, 506921, 0.00, 69.46
2012-11-14 19:32:33, 0, 507088, 0.00, 69.46
2012-11-14 19:32:38, 0, 506895, 0.00, 69.46
2012-11-14 19:32:43, 0, 507020, 0.00, 69.46
2012-11-14 19:32:48, 0, 506976, 0.00, 69.46

```

Code Example 9: upacNIProvPerf Statistics File Sample**6.7.3 upacNIProvPerf Console Output Sample**

```

020: UpdRate:    101396, CPU:    0.00%, Mem:   69.46MB
025: UpdRate:    101384, CPU:    0.00%, Mem:   69.46MB
030: UpdRate:    101417, CPU:    0.00%, Mem:   69.46MB
035: UpdRate:    101379, CPU:    0.00%, Mem:   69.46MB
040: UpdRate:    101404, CPU:    0.00%, Mem:   69.46MB
045: UpdRate:    101395, CPU:    0.00%, Mem:   69.46MB

```

Code Example 10: upacNIProvPerf Console Output Sample

Chapter 7 upacTransportPerf

7.1 Overview

upacTransportPerf measures the performance of the various API transport layers. **upacTransportPerf** does not use OMM messages, instead sending opaque content with minimal encoding and decoding. To enforce the proper ordering of data, **upacTransportPerf** embeds a sequence number in each buffer.

upacTransportPerf can act as a server or client. A typical use case is to start **upacTransportPerf** as a server and then start **upacTransportPerf** as a client to connect to the server.

7.2 Threading and Scaling

The Transport API allows calls from multiple threads, such that applications can scale their work across multiple cores (via multi-threading) to handle multiple connections. To support multi-threading, each application must enable global locking when calling **rsslInitialize** to ensure that shared resources in the Transport API are protected.

You can configure **upacTransportPerf** for multiple threads by using the **-threads** option. The result depends on whether the application is run as a server or client.

- When running as a server, each thread is used to balance connections (similar to **upacProvPerf**).
- When running as a client, each thread creates its own connection (similar to **upacConsPerf**).

The main thread monitors the other threads and collects and reports their statistics.

7.3 upacTransportPerf Life Cycle

The lifecycle of **upacTransportPerf** is divided into the following phases:




Figure 16. upacTransportPerf Lifecycle 1

1. Application and Transport API initialization.

upacTransportPerf loads its configuration, initializes the Transport API, and starts the transport thread(s) that read and write data. The main thread now:

- If acting as server, accepts client connections and passes them to one of the transport threads.
- If acting as client, connects out to server.
- Periodically collects and writes statistics from the transport thread(s) for the remainder of the test.

2. Publishing and receiving opaque content.

The transport thread performs the following actions until its run time expires:

- Adds new connections passed from the main thread.
- Sends bursts of messages on existing connections, as specified by command line arguments.
- Uses remaining time segments to read from the transport and process incoming data.

3. Application shutdown and cleanup.

The transport thread stops. The main thread collects any remaining data from the transport threads, cleans them up, and writes the final summary statistics. Finally, the main thread uninitialized the Transport API, cleans up remaining resources, and exits.




Figure 17. upacTransportPerf Application Flow

7.4 Message Payload

upacTransportPerf writes a message of the configured size obeying the following rules:

- The message starts with an 8-byte sequence number, which is checked by the receiver of the message.
- If the message contains a timestamp, the stamp is added after the sequence number as an 8-byte integer.
- The remainder of the message is set to zeros.

7.5 Latency Measurement

upacTransportPerf writes a timestamp (with nanosecond granularity) as part of the message payload (as described in Section 7.4). To determine latency, the receiving **upacTransportPerf** reads the timestamp and compares it to the current time.

► Transport Performance Latency Measurement Sequence:

1. Get the current time (**t1**).
2. Obtain an output buffer using **rsslGetBuffer**.
3. Encode the message, including time **t1**.
4. Pass the message to the API, which passes it to underlying transport.

To determine latency, the consuming application reads the timestamp from the payload and compares it against the current time.

7.6 upacTransportPerf Configuration Options

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-appType	server	Configures the type of application that upacTransportPerf simulates (available settings include server or client).
-busyRead	<none>	Configures the application to continually read rather than use notifications.
-compressionLevel	5	When upacTransportPerf is configured to compress messages, -compressionLevel sets the level of compression.
-compressionType	<none>	Sets the type of compression upacTransportPerf uses when compressing messages. By default, upacTransportPerf does not compress its messages. Example: zlib
-connType	socket	Specifies the connection type that the application uses. Supported options are socket , http , encrypted , reliableMCast , and shmem .
-directWrite	False	Sets whether the RSSL_WRITE_DIRECT_SOCKET_WRITE flag is used to call rsslWrite . Using this flag causes the write to attempt to bypass the Transport API output queue, reducing latency but decreasing throughput capability.

Table 5: upacTransportPerf Configuration Options


COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-h	localhost	When using TCP socket or shared memory connection types, specifies the hostname to which the client application connects.
-highWaterMark	6000	Configures the Transport API's "High-water Mark," the amount of data (in bytes) that the Transport API queues before flushing it to the network. Adjusting this might provide a trade-off of throughput vs. latency when writing large bursts of data.
-if	<none>	Configures interfaceName (an RsslConnectOptions/RsslBindOptions parameter), which configures the NIC that the provider uses for its server. On computers connected to multiple networks, you can use this parameter to force the provider to use a specific network.
-inputBufs	50	Configures numInputBuffers (an RsslConnectOptions/RsslBindOptions parameter), which sets the size of the Transport API's input queue. Use a value large enough to accommodate incoming data and minimize network read operations.
-latencyFile	<none>	Configures upacTransportPerf to log the latency of data from updates and posts, and specifies the name of the file in which upacTransportPerf stores the results.
-latencyMsgRate	10	Sets the number of messages that upacTransportPerf sends per second to measure latency. This cannot be larger than -msgRate .
-maxFragmentSize	6144	Configures maxFragmentSize (an RsslBindOptions parameter), which controls the size (in bytes) of the buffers in the Transport API's buffer pool.
-mcastStats	(no argument)	Enables upacTransportPerf to print multicast statistics. To print multicast statistics, upacTransportPerf must also enable channel locks in the Transport API.
-msgRate	100000	Sets the number of messages that upacTransportPerf sends per second. -msgRate cannot be less than -latencyMsgRate as latency messages are a type of message and hence included in this number. For example: using -latencyMsgRate 1000 -msgRate 1000 sends 1000 messages per second with each message carrying latency data.
-msgSize	76	Configures the size (in bytes) of messages sent by upacTransportPerf .
-noDisplayStats	(no argument)	Turns off printing statistics to the screen.
-outputBufs	5000	Configures guaranteedOutputBuffers (an RsslConnectOptions/RsslBindOptions parameter), which sets the minimum guaranteed number of output buffers created for each channel.  Tip: Set this value large enough so that the provider does not run out of buffers while writing but low enough so as not to waste memory and slow down performance.
-p	14002	When using TCP socket or shared memory connection types, specifies the port number over which the consumer connects. This option applies to both types of applications (server or client).

Table 5: upacTransportPerf Configuration Options (Continued)

COMMAND-LINE OPTION	DEFAULT	DESCRIPTION
-pack	No packing	Sets the number of messages to pack in each buffer using <code>rsslPackBuffer</code> .
-ra	<none>	When using the multicast connection type, configures the multicast receive address.
-recvBufSize	<none>	Configures the system receive buffer size. By default, the OS setting is used.
-reflectMsgs	False	Sets upacTransportPerf to reflect back all received messages.
-rp	<none>	When using the multicast connection type, configures the multicast receive port.
-runTime	300	Sets how long upacTransportPerf runs, in seconds.
-sa	<none>	When using the multicast connection type, configures the multicast send address.
-sendBufSize	<none>	Configures the system send buffer size. By default, the OS setting is used.
-sp	<none>	When using the multicast connection type, configures the multicast send port.
-statsFile	TransportStats	Specifies the base name used in writing the application's test statistics.
-summaryFile	TransportSummary.out	Specifies the name of the file to which upacTransportPerf writes the application's test summary.
-tcpDelay	False	Configures tcp_nodelay (an RsslConnectOptions/RsslBindOptions parameter) which sets whether the underlying connection uses Nagle's algorithm; a method for more efficiently using network packet headers by batching data. By default, upacTransportPerf disables Nagle's algorithm due to the increased latency from buffering. <code>-tcpDelay</code> enables it, which can raise throughput capability.
-threads	<none>	Sets the number of threads that the application creates and the CPU core to which each thread binds.
-tickRate	1000	Sets the number of cycles (i.e., ticks) per second executed by the upacTransportPerf main loop. Adjusting the tick rate changes the size of request/post bursts: a higher tick rate results in smaller individual bursts and smoother traffic.
-u	<none>	When using the multicast connection type, configures the unicast port.
-writeStatsInterval	5	Configures how often (in seconds) upacTransportPerf prints stats to screen and to the statistics file.

Table 5: upacTransportPerf Configuration Options (Continued)

7.7 Input

upacTransportPerf does not require input files.

7.8 Output

During the test, **upacTransportPerf** records the following statistics:

- Message rate (sent and received)
- Data rate (sent and received)
- Latency statistics
- CPU and memory usage

Note: The data sent rate is taken from **rsslWrite**'s **bytesWritten** parameter. This value includes any internal headers used in transport, packing, and/or compression. The data received rate is based on message lengths returned from **rsslRead**, which does not indicate whether data was compressed and does not include any transport header overhead. As a result, rates may differ between the sender and receiver.

For more detailed output file information, refer to Chapter 10.

7.8.1 upacTransportPerf Summary File Sample

```

--- TEST INPUTS ---

        Runtime: 180 sec
Connection Type: socket
        Hostname: localhost
           Port: 14002
        App Type: client
    Thread List: 7
        Busy Read: No
           Msg Size: 76
        Msg Rate: 100000
Latency Msg Rate: 10
    Output Buffers: 5000
Max Fragment Size: 6144
    Send Buffer Size: 0(use default)
    Recv Buffer Size: 0(use default)
    High Water Mark: 0(use default)
    Compression Type: none(0)
Compression Level: 5
    Interface Name: (use default)
        Tcp_NoDelay: Yes
           Tick Rate: 1000
Use Direct Writes: No
    Latency Log File: (none)
        Summary File: TransportSummary_24806.out
           Stats File: TransportStats

```

```

Write Stats Interval: 5
    Display Stats: Yes
    Packing: No

--- OVERALL SUMMARY ---

Statistics:
  Latency avg (usec): 19.896
  Latency std dev (usec): 160.616
  Latency max (usec): 5304.720
  Latency min (usec): 9.052
  Sampling duration(sec): 179.98
  Msgs Sent: 17998200
  Msgs Received: 17999300
  Data Sent (MB): 1355.99
  Data Received (MB): 1304.58
  Avg. Msg Sent Rate: 100000
  Avg. Msg Recv Rate: 100006
  Avg. Data Sent Rate (MB): 7.53
  Avg. Data Recv Rate (MB): 7.25
  CPU/Memory samples: 36
  CPU Usage max (%): 0.18
  CPU Usage min (%): 0.00
  CPU Usage avg (%): 0.00
  Memory Usage max (MB): 31.20
  Memory Usage min (MB): 21.93
  Memory Usage avg (MB): 28.90
  Process ID: 24806

```

Code Example 11: upacTransportPerf Summary File Sample

7.8.2 upacTransportPerf Statistics File Sample

```

UTC, Msgs sent, Bytes sent, Msgs received, Bytes received, Latency msgs received, Latency avg (usec),
  Latency std dev (usec), Latency max (usec), Latency min (usec), CPU usage (%), Memory (MB)
2012-11-14 06:35:59, 498000, 39342000, 499100, 37931600, 50, 206.135, 954.202, 5304.720, 9.902, 0.18,
  21.93
2012-11-14 06:36:04, 500000, 39500000, 500000, 38000000, 50, 14.107, 2.554, 19.491, 9.227, 0.00, 21.99
2012-11-14 06:36:09, 500000, 39500000, 500000, 38000000, 50, 14.836, 2.667, 19.105, 10.180, 0.00, 21.99
2012-11-14 06:36:14, 500000, 39500000, 500000, 38000000, 50, 14.720, 2.710, 19.046, 9.721, 0.00, 21.99
2012-11-14 06:36:19, 500000, 39500000, 500000, 38000000, 50, 14.434, 2.293, 19.497, 10.722, 0.00, 21.99
2012-11-14 06:36:24, 500000, 39500000, 500000, 38000000, 50, 14.133, 2.401, 19.122, 9.683, 0.00, 21.99
2012-11-14 06:36:29, 500000, 39500000, 500000, 38000000, 50, 14.878, 2.527, 19.159, 9.773, 0.00, 21.99
2012-11-14 06:36:34, 500000, 39500000, 500000, 38000000, 50, 14.874, 2.719, 20.240, 10.386, 0.00, 23.96
2012-11-14 06:36:39, 500000, 39500000, 500000, 38000000, 50, 14.460, 2.303, 19.820, 9.975, 0.00, 23.96

```

Code Example 12: upacTransportPerf Statistics File Sample

7.8.3 upacTransportPerf Console Output Sample

```

030:
Sent: MsgRate: 100000, DataRate: 7.534MBps
Recv: MsgRate: 100000, DataRate: 7.248MBps
Latency (usec): Avg: 14.133 StdDev: 2.401 Max: 19.122 Min: 9.683, Msgs: 50
CPU: 0.00% Mem: 21.99MB
035:
Sent: MsgRate: 100000, DataRate: 7.534MBps
Recv: MsgRate: 100000, DataRate: 7.248MBps
Latency (usec): Avg: 14.878 StdDev: 2.527 Max: 19.159 Min: 9.773, Msgs: 50
CPU: 0.00% Mem: 21.99MB
040:
Sent: MsgRate: 100000, DataRate: 7.534MBps
Recv: MsgRate: 100000, DataRate: 7.248MBps
Latency (usec): Avg: 14.874 StdDev: 2.719 Max: 20.240 Min: 10.386, Msgs: 50
CPU: 0.00% Mem: 23.96MB
045:
Sent: MsgRate: 100000, DataRate: 7.534MBps
Recv: MsgRate: 100000, DataRate: 7.248MBps
Latency (usec): Avg: 14.460 StdDev: 2.303 Max: 19.820 Min: 9.975, Msgs: 50
CPU: 0.00% Mem: 23.96MB
050:
Sent: MsgRate: 100000, DataRate: 7.534MBps
Recv: MsgRate: 100000, DataRate: 7.248MBps
Latency (usec): Avg: 14.213 StdDev: 2.291 Max: 19.355 Min: 10.093, Msgs: 50
CPU: 0.00% Mem: 29.23MB

```

Code Example 13: upacTransportPerf Console Output Sample

Chapter 8 Performance Measurement Scenarios

8.1 Interactive Provider to Consumer, Through TREP

You can measure interactive providers by connecting **upacConsPerf** to an ADS, the ADS to an ADH,¹ and finally the ADH with an instance of **upacProvPerf**. You can perform this test with caching enabled or disabled in the ADH or ADS, as **upacProvPerf** acts as the cache of record in this scenario.




Figure 18. Interactive Provider to Consumer on TREP

► To run a basic performance measurement:

To run a basic performance measurement using a setup like the one pictured in Figure 18, run **upacProvPerf** and **upacConsPerf** with the following commands:²

```
upacProvPerf -p 17008 -serviceName TEST_FEED

upacConsPerf -h adshost -p 14002 -serviceName TEST_FEED
```

1. Via the RRCP backbone.

2. The options on these command lines assume TEST_FEED is the service being used and 17008 is the port number. Modify these example values as necessary.

8.2 Interactive Provider to Consumer, Direct Connect

You can measure the interactive providers of data by connecting **upacConsPerf** directly to **upacProvPerf**.




Figure 19. Interactive Provider to Consumer, Direct Connect

► To run a basic performance measurement:

Using their default configuration options, you can run this test without any additional command-line options. Simply run the provider and consumer applications as follows:

```
upacProvPerf
```

```
upacConsPerf
```

8.3 Non-Interactive Provider to Consumer, Through TREP

You can measure non-interactive providers on TREP by connecting **upacConsPerf** to an ADS, the ADS with an ADH,³ and finally the **upacNIProvPerf** to the ADH. The ADH must have caching enabled, because it acts as the cache of record in this scenario.

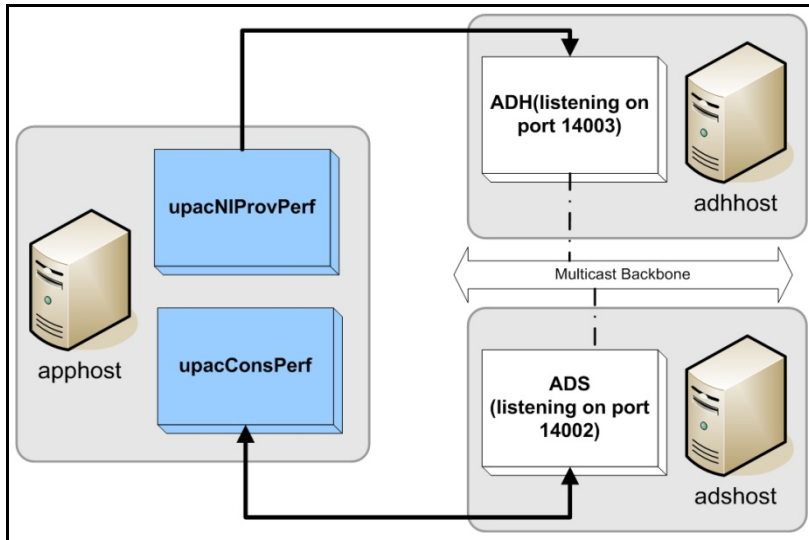


Figure 20. NIProv to Consumer on the TREP

upacConsPerf may receive a Closed status if it requests an item not yet provided by **upacNIProvPerf** to the ADH cache. To ensure the test completes successfully, you must do either one of the following:

1. Preload the ADH cache. **upacNIProvPerf** must have provided refreshes for all of its items to the ADH before **upacConsPerf** connects to the ADS.
2. Configure the ADH to provide temporary refreshes in place of the uncached items. **upacConsPerf** knows to allow these images, and does not count them towards the image retrieval time, due to their Suspect data state.

For more details on this configuration, refer to the *ADH Software Installation Manual*.

► To run a basic performance measurement:

To run a basic performance measurement using a setup like the one pictured in Figure 20, run **upacNIProvPerf** and **upacConsPerf** with the following command-line options:⁴

```
upacNIProvPerf -h adhhost -p 14003 -serviceName TEST_FEED

upacConsPerf -h adshost -p 14002 -serviceName TEST_FEED
```

3. Via the RRCP backbone.

4. These options assume the provided service is TEST_FEED. Modify the example's values as necessary.

8.4 Consumer Posting on the TREP

To measure posting performance on the TREP, connect the **upacConsPerf** to an ADS, the ADS to an ADH,⁵ and finally the **upacNIProvPerf** to the ADH. The ADH must have caching enabled, because it acts as the cache of record in this scenario. As the posted messages return from the TREP, the consumer can distinguish them via the presence of their **RssIPostUserInfo**. When configured to do so, **upacConsPerf** embeds timestamps in some of its posts which it uses to measure round-trip latency.

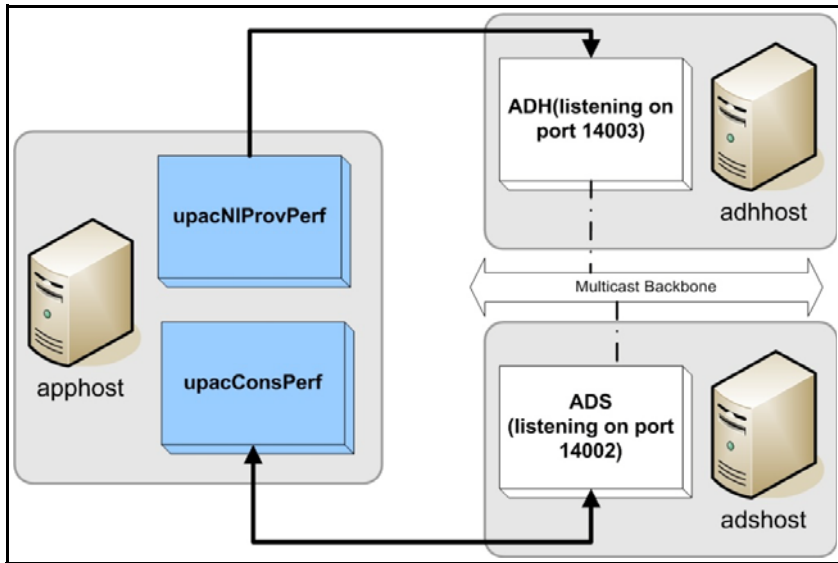


Figure 21. Consumer Posting to TREP

Update traffic is optional. If you want to test posting without updates, configure **upacNIProvPerf** by specifying `-updateRate 0` `-latencyUpdateRate 0` in the command line.

Additionally, if you want only posting traffic, you do not need to run a provider application. You can configure the TREP to provide the necessary service information and refresh content. For more details on this configuration, refer to the *ADH Software Installation Manual*.

► To run a basic performance measurement:

To run a basic performance measurement using a setup like the one pictured in Figure 21, run **upacNIProvPerf** and **upacConsPerf** as follows:⁶

```
upacNIProvPerf -h adhost -p 14003 -serviceName TEST_FEED

upacConsPerf -h adhost -p 14002 -serviceName TEST_FEED -postingRate 10000 -postingLatencyRate 10
```

5. Via the RRCP backbone.

6. These options assume TEST_FEED is the service being provided. Modify the example values as necessary.

8.5 Transport Performance, Direct Connect with TCP

You can measure the performance of the transport⁷ in a direct connect scenario. This test can determine the throughput and latency you can achieve using different transport types supported in the Transport API. This scenario shows the TCP Socket transport type.




Figure 22. Transport Performance, TCP Direct Connect

► To run a basic performance measurement:

To run a basic performance measurement using the setup pictured in Figure 22, run two instances of **upacTransportPerf** as follows:

```
upacTransportPerf -p 15000 -runTime 900 -connType socket -msgRate 100000 -latencyMsgRate 1000
                  -tickRate 1000 -appType server

upacTransportPerf -p 15000 -connType socket -runTime 60 -latencyMsgRate 0 -msgRate 0 -tickRate 1000
                  -appType client
```

7. Other than timestamp and sequence number encoding and decoding, this application does not perform other content operations or inspections.

8.6 Transport Performance, Direct Connect with TCP, Reflection

To measure latency across hosts, **upacTransportPerf** supports reflecting messages. When reflecting is enabled, instead of **upacTransportPerf** producing its own messages, each read message is written back as-is to the same connection. Reflecting is enabled by the **-reflectMsgs** command-line option and may be done by servers or clients.




Figure 23. Transport Performance, TCP Direct Connect with Reflection

► To run a Basic Performance Measurement:

To run a basic performance measurement using a setup like the one pictured in Figure 23, run two instances of **upacTransportPerf** with the following command-line options:

```
upacTransportPerf -p 15000 -runTime 900 -connType socket -tickRate 1000 -appType server -reflectMsgs
upacTransportPerf -h apphost1 -p 15000 -connType socket -runTime 60 -tickRate 1000 appType client
    -msgRate 100000 -latencyMsgRate 1000
```

8.7 Transport Performance, Direct Connect with Multicast

Because of the nature of multicast, both applications are configured as clients. Also, this example uses separate 'send' and 'receive' networks.




Figure 24. Transport Performance, Multicast Direct Connect

► To run a basic performance measurement:

To run a basic performance measurement using a setup like the one pictured in Figure 24, run two instances of **upacTransportPerf** as follows:

```
upacTransportPerf -runTime 90 -connType reliableMCast -u 12345 -sp 15001 -sa 235.1.1.3 -rp 15002
                  -ra 235.1.1.4 -latencyMsgRate 1000 -msgRate 10000 -tickRate 1000 -appType client

upacTransportPerf -u 14006 -sp 15002 -sa 235.1.1.4 -rp 15001 -ra 235.1.1.3 -connType reliableMCast
                  -runTime 60 -appType client -msgRate 0 -latencyMsgRate 0 -tickRate 1000
```

8.8 Transport Performance, Direct Connect with Shared Memory

The following example uses a small **maxFragmentSize** to reduce the size of the shared memory segment and uses the **-threads** option. The provider and consumer each use one thread: the provider thread is bound to core 0, and the consumer thread is bound to core 1. Such a setup presumes the system has at least two cores.




Figure 25. TransportPerf, Shared Memory Direct Connect

► To run a basic performance measurement:

To run a basic performance measurement using a setup like the one pictured in Figure 25, run two instances of **upacTransportPerf** as follows:

```
upacTransportPerf -p 15000 -runTime 90 -connType shmem -msgRate 10000 -latencyMsgRate 1000
                  -tickRate 1000 -appType server -outputBufs 9000 -maxFragmentSize 256 -threads 0

upacTransportPerf -p 15000 -connType shmem -runTime 60 -latencyMsgRate 0 -msgRate 0 -tickRate 1000
                  -appType client -threads 1
```

Chapter 9 Input File Details

9.1 Message Content File and Format

The message data XML file (**MsgData.xml**) provided with the Performance Suite describes sample data for the refreshes, updates, and posts encoded by the tools. You can customize **MsgData.xml** to suit desired test scenarios.

The XML file must contain data for:

- One refresh message.
- At least one update message.
- At least one post message, if posting from **upacConsPerf**.
- At least one generic message, if configured for exchanging generic messages.

Refresh data provides the image for each item provided by **upacProvPerf** or **upacNIProvPerf**. When providing updates, provider tools encode update messages in a round-robin manner for each item. Likewise, when posting, the **upacConsPerf** encodes posts in a round-robin fashion for each item.

9.1.1 Encoding Fields

Performance tools can encode in their fields any of the primitive types supported by the Transport API.

For types such as **RsslReal** and **RsslDate**, the input follows the conversion format of the Transport API's string conversion function (e.g. **rsslNumericStringToReal**, **rsslDateStringToDate**).

Each field must have the correct type for its ID according to the dictionary loaded by the tool. Fields are validated by the message data parser.

9.1.2 Sample Update Message

```
<updateMsg>
  <dataBody>
    <fieldList>
      <fieldEntry fieldId="1025" dataType="RSSL_DT_TIME" data=" 14:40:32:000"/>
      <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="52832000"/>
      <fieldEntry fieldId="115" dataType="RSSL_DT_ENUM" data="2"/>
      <fieldEntry fieldId="1000" dataType="RSSL_DT_RMTES_STRING" data="-"/>
      <fieldEntry fieldId="22" dataType="RSSL_DT_REAL" data="401.50"/>
      <fieldEntry fieldId="114" dataType="RSSL_DT_REAL" data="3.49"/>
      <fieldEntry fieldId="30" dataType="RSSL_DT_REAL" data="18"/>
      <fieldEntry fieldId="25" dataType="RSSL_DT_REAL" data="401.54"/>
      <fieldEntry fieldId="31" dataType="RSSL_DT_REAL" data="10"/>
      <fieldEntry fieldId="293" dataType="RSSL_DT_RMTES_STRING" data="NAS"/>
      <fieldEntry fieldId="3298" dataType="RSSL_DT_ENUM" data="43"/>
      <fieldEntry fieldId="296" dataType="RSSL_DT_RMTES_STRING" data="NAS"/>
      <fieldEntry fieldId="3297" dataType="RSSL_DT_ENUM" data="43"/>
    </fieldList>
  </dataBody>
</updateMsg>
```

Code Example 14: Sample Update Message

9.1.3 Sample MarketByOrder Data

Performance tools also support MarketByOrder data, however it is currently experimental. To allow tools to provide MarketByOrder data, you can add the following data to **MsgData.xml**:

```
<!-- MarketByOrder -->
<marketByOrderMsgList>

<!-- ORDER_SIDE enumerations: BID is 1, ASK is 2 -->

<refreshMsg>
  <dataBody>
    <map>
      <fieldSetDefs>
        <fieldSetDef setId="0">
          <fieldSetDefEntry fieldId="3427" dataType="RSSL_DT_REAL" />
          <fieldSetDefEntry fieldId="3429" dataType="RSSL_DT_REAL" />
          <fieldSetDefEntry fieldId="3855" dataType="RSSL_DT_UINT_2" />
          <fieldSetDefEntry fieldId="212" dataType="RSSL_DT_RMTES_STRING" />
          <fieldSetDefEntry fieldId="3428" dataType="RSSL_DT_ENUM" />
        </fieldSetDef>
      </fieldSetDefs>
      <mapEntry action="RSSL_MPEA_ADD_ENTRY" key="100" >
```

```

        <fieldList setId="0">
            <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189639.67"/>
            <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018963967"/>
            <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="0"/>
            <fieldEntry fieldId="212" dataType="RSSL_DT_RMTES_STRING" data="MarketMaker1"/>
            <fieldEntry fieldId="3428" dataType="RSSL_DT_ENUM" data="1"/>
        </fieldList>
    </mapEntry>
    <mapEntry action="RSSL_MPEA_ADD_ENTRY" key="101" >
        <fieldList setId="0">
            <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189638.67"/>
            <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018963467"/>
            <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="500"/>
            <fieldEntry fieldId="212" dataType="RSSL_DT_RMTES_STRING" data="MarketMaker2"/>
            <fieldEntry fieldId="3428" dataType="RSSL_DT_ENUM" data="1"/>
        </fieldList>
    </mapEntry>
    <mapEntry action="RSSL_MPEA_ADD_ENTRY" key="102" >
        <fieldList setId="0">
            <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189637.67"/>
            <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018962967"/>
            <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="1000"/>
            <fieldEntry fieldId="212" dataType="RSSL_DT_RMTES_STRING" data="MarketMaker3"/>
            <fieldEntry fieldId="3428" dataType="RSSL_DT_ENUM" data="2"/>
        </fieldList>
    </mapEntry>
    <mapEntry action="RSSL_MPEA_ADD_ENTRY" key="103" >
        <fieldList setId="0">
            <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189636.67"/>
            <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018962467"/>
            <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="1500"/>
            <fieldEntry fieldId="212" dataType="RSSL_DT_RMTES_STRING" data="MarketMaker4"/>
            <fieldEntry fieldId="3428" dataType="RSSL_DT_ENUM" data="2"/>
        </fieldList>
    </mapEntry>
    <mapEntry action="RSSL_MPEA_ADD_ENTRY" key="104" >
        <fieldList setId="0">
            <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189635.67"/>
            <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018961967"/>
            <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="2000"/>
            <fieldEntry fieldId="212" dataType="RSSL_DT_RMTES_STRING" data="MarketMaker5"/>
            <fieldEntry fieldId="3428" dataType="RSSL_DT_ENUM" data="2"/>
        </fieldList>
    </mapEntry>
</map>
</dataBody>
</refreshMsg>

<updateMsg>
    <dataBody>

```

```

<map>
  <mapEntry action="RSSL_MPEA_UPDATE_ENTRY" key="100" >
    <fieldList>
      <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189638.67"/>
      <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018963867"/>
      <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="0"/>
    </fieldList>
  </mapEntry>
  <mapEntry action="RSSL_MPEA_UPDATE_ENTRY" key="101" >
    <fieldList>
      <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189635.67"/>
      <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018963567"/>
      <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="500"/>
    </fieldList>
  </mapEntry>
  <mapEntry action="RSSL_MPEA_UPDATE_ENTRY" key="102" >
    <fieldList>
      <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189632.67"/>
      <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018963267"/>
      <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="1000"/>
    </fieldList>
  </mapEntry>
  <mapEntry action="RSSL_MPEA_UPDATE_ENTRY" key="103" >
    <fieldList>
      <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189629.67"/>
      <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018962967"/>
      <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="1500"/>
    </fieldList>
  </mapEntry>
  <mapEntry action="RSSL_MPEA_UPDATE_ENTRY" key="104" >
    <fieldList>
      <fieldEntry fieldId="3427" dataType="RSSL_DT_REAL" data="360287970189626.67"/>
      <fieldEntry fieldId="3429" dataType="RSSL_DT_REAL" data="36028797018962667"/>
      <fieldEntry fieldId="3855" dataType="RSSL_DT_UINT" data="2000"/>
    </fieldList>
  </mapEntry>
</map>
</dataBody>
</updateMsg>

</marketByOrderMsgList>

```

Code Example 15: Sample MarketByOrder Data

9.2 Item List File

The Item List File configures the full list of items as requested by **upacConsPerf** or published by **upacNIProvPerf**. Each entry specifies the item's name and how it is requested. The file must contain enough entries to satisfy the number of items needed by the respective tool.

The sample file **350k.xml** contains 350,000 items, some of which allow posting.

9.2.1 Item Attributes

ATTRIBUTE NAME	DEFAULT	DESCRIPTION
domain	(none, required)	Specifies the domain from which the item is requested. This must be set to MarketPrice .
genMsg	"false"	If set to true, generic messages are sent for this item (if generic messages are enabled).
name	(none, required)	Specifies the name used in the MsgKey when requesting the item.
post	"false"	If set to true , upacConsPerf sends posts to this item (if posting is enabled).
snapshot	"false"	If set to true , upacConsPerf requests this item as a snapshot (i.e., without setting the RSSL_RQF_STREAMING flag on the request).

Table 6: Item Attributes

9.2.2 Sample Item List File

```
<itemList>
  <item domain="MarketPrice" name="RDT1" post="true" genMsg="true" />
  <item domain="MarketPrice" name="RDT2" post="true" />
  <item domain="MarketPrice" name="RDT3" post="true" />
  <item domain="MarketPrice" name="RDT4" post="true" />
  <item domain="MarketPrice" name="RDT5" post="true" />
  <item domain="MarketPrice" name="RDT6" post="true" />
  <item domain="MarketPrice" name="RDT7" post="true" />
  <item domain="MarketPrice" name="RDT8" />
  <item domain="MarketPrice" name="RDT9" />
  <item domain="MarketPrice" name="RDT10" />
  <item domain="MarketPrice" name="RDT11" />
  <item domain="MarketPrice" name="RDT12" />
  <item domain="MarketPrice" name="RDT13" />
  <item domain="MarketPrice" name="RDT14" />
  <item domain="MarketPrice" name="RDT15" />
  <item domain="MarketPrice" name="RDT16" />
  <item domain="MarketPrice" name="RDT17" />
  <item domain="MarketPrice" name="RDT18" />
</itemList>
```

Code Example 16: Sample Item List File

Chapter 10 Output File Details

10.1 Overview

Applications in the Performance Suite send similar output to the console and to files. Each application can configure its output using the configuration parameters:

- **writeIntervalStats** (1 to *n*): The interval (in seconds) at which timed statistics are written to files.
- **noDisplayStats**: Disables statistics output to the console.

Providers and consumers output different statistics but in a similar fashion. Each application can be configured to output a summary file, a statistics file, and in the case of the consumer, a latency file comprised of individual latencies for each received latency item.

10.2 Output Files and Their Descriptions

You can configure the names of output files, though applications append the client number to their stats and latency files. So for example, a horizontal scaling test with two consumer threads produces two statistics files: **ConsStats1.csv** and **ConsStats2.csv**.

Default output filenames (and the associated parameters you use to generate the files) are as follows:

PARAMETER	DEFAULT	DESCRIPTION
-latencyFile	(none)	Specifies the filename of the latency file produced.
-statsFile	<i>ToolType</i> Statsclient.csv ^a	Specifies the filename of the statistics file produced.
-summaryFile	<i>ToolType</i> Summary.txt	Specifies the filename of the summary file produced.

Table 7: Performance Suite Applications and Associated Configuration Files

a. Where *ToolType* is either **Cons**, **Prov**, **NIProv**, or **Transport**.

10.3 Latency File

The latency file is a comma-separated value file containing individual latencies, in microseconds, for timestamps received during the test. It is only created by **upacConsPerf** and **upacTransportPerf**.

Note: Due to the potentially large amount of output in scenarios that use a high latency message rate, this file is not produced by default.

The interval in seconds that statistics are written to the file is controlled by the **writeStatsInterval** configuration parameter, which defaults to **5**.

```
Message type, Send time, Receive time, Latency (usec)
Upd, 353725032296, 353725032521, 225
Upd, 353725045319, 353725045569, 250
Upd, 353725092300, 353725092521, 221
Pst, 353724892323, 353724894740, 2417
Pst, 353724925257, 353724926441, 1184
Pst, 353725105324, 353725106762, 1438
Upd, 353725359645, 353725359859, 214
Upd, 353725610354, 353725610619, 265
```

Code Example 17: Sample ConsLatency.csv Showing Update and Post latencies during a Test Run

10.4 File Import

You can import output **.csv** files into data analysis software. For example, you can use Microsoft Excel and Microsoft Access to import and quickly analyze your test results. Shown below are graphs created in Excel after importing a statistics **.csv** file for a test run. Note that these are sample graphs and do not imply the real performance results of the tool suite.




Figure 26. Sample Excel Graph from ConsStats1.csv




Figure 27. Sample Excel Graph of Latencies Over a 15-second Steady State Interval from ConsLatency1.csv

Chapter 11 Performance Best Practices

11.1 Overview

The Performance Test Tools Suite leverages a number of features of the Transport API to achieve high throughput and low latency when sending and receiving messages. This section briefly describes test tool features, the features' benefits, and how the tools use them. For more details on each feature, refer to the *Transport API Developers Guide*.

11.2 Transport Best Practices

11.2.1 `rsslRead`

When calling `rsslRead`, the `readret` parameter indicates whether more content is available for immediate processing. Because this content might have already been read from the underlying transport, an I/O notifier might not inform an application that this content is available. When the value is greater than **RSSL_RET_SUCCESS (0)**, the application should call `rsslRead` again without waiting for further notification. Even if no message was returned, `rsslRead` can return a value greater than **RSSL_RET_SUCCESS**.

All applications follow a reading model that leverages the Transport API's queuing and notification mechanisms to read quickly without making unnecessary function calls:

1. Wait for notification on the descriptors associated with a channel using its preferred I/O notifier. The application will only call `rsslRead` on channels that indicate the presence of readable data.
2. Keep calling `rsslRead` as long as the return value is greater than **RSSL_RET_SUCCESS**. Such return values indicate that there is still data to process in the Transport API's input buffer.

11.2.2 `rsslWrite`, `rsslFlush`

To make efficient use of underlying transport function calls, the `rsslWrite` function passes messages to an outbound queue of the specified priority, rather than immediately writing the message to the network. `rsslWrite` indicates that there is queued content by returning a value greater than **RSSL_RET_SUCCESS**.

The network write occurs if:

- The application calls `rsslFlush`.
- The **RSSL_WRITE_DIRECT_SOCKET_WRITE** flag is passed into the `rsslWrite` function.
- The amount of queued data reaches a configurable highwater mark (i.e., using the `-highWaterMark` option), which causes `rsslWrite` to pass queued content to the underlying transport.

You can use a simple algorithm to write large amounts of content efficiently while still maintaining low latency:

- Write all currently desired content via `rsslWrite`, relying on the highwater mark to periodically flush.
- When there is no more content to write, call `rsslFlush` to flush any remaining data. After all data is written to the network, `rsslFlush` returns **RSSL_RET_SUCCESS**.

Note: A positive value returned from `rsslWrite()` means there is queued content. It is not passed to the underlying transport unless the application performs one of the previously mentioned actions.

Transport API Performance Tools observe the following model when writing message bursts:

1. Write the entire burst of messages using `rsSlWrite`.
2. After the burst finishes, check the return value from the last call to `rsSlWrite`. If the value is greater than `RSSL_RET_SUCCESS`, data remains in the Transport API's output queue. Use the channel's file descriptor in conjunction with an I/O notifier to notify the application when it can flush remaining data.
3. Call `rsSlFlush` on channels indicated by the notifier. Continue to invoke the notifier and `rsSlFlush` until all data is successfully provided to the underlying transport.

Each application has a command-line option (`-highWaterMark`) for adjusting the high-water mark.

11.2.3 Packing

To efficiently use buffer space, reduce the number of writes to the transport,¹ and improve throughput, applications can use the `rsSlPackBuffer` function to pack multiple messages into a single buffer. Each call to `rsSlPackBuffer` adds another message to the existing buffer. Though packing messages can help increase throughput, overall effects will vary depending on the type of transport you use. In general:

- Slower transports benefit more than faster transports
- More saturated transports benefit more than less-saturated ones.

Because packing adds latency, it is important that you weigh the trade-offs when deciding to employ packing. For example, if 15 messages are packed into a single buffer, the first message will sit in the buffer longer before being written than the last message added to the buffer. An application can reduce this latency by sending the packed buffer before it is full, often through the use of a timing mechanism that indicates a latency threshold has been reached. In general, the more messages packed into a buffer, the higher the latency penalty.

`upacProvPerf`, `upacNIProvPerf`, and `upacTransportPerf` perform the following steps when configured to pack buffers:

1. Get a packable buffer of a specified size, using `rsSlGetBuffer`.
2. Encode messages into the buffer, calling `rsSlPackBuffer` after each message is encoded.
3. After the configured number of messages are packed or when the buffer is full, `rsSlWrite` is called to pass the packed buffer to the Transport API.

11.2.4 High-water Mark

Higher throughput is usually achieved by making a small number of large writes to the transport instead of doing a large number of small writes. For example, writing one 6000-byte buffer is generally more efficient than writing 1000 six-byte buffers. To achieve higher efficiencies, the Transport API employs the concept of a high-water mark. When the application calls `rsSlWrite`, the Transport API does not always immediately pass the buffer to the transport; instead, the Transport API passes data to the transport after the size of its buffer reaches the high-water mark.² Note that the high-water mark applies to written data and not to the number of bytes requested in the call to `rsSlGetBuffer`.

For example, assume a high-water mark of 6144 bytes. If an application gets a buffer, encodes 500 bytes of content, and passes this to `rsSlWrite`, the high-water mark will be triggered after thirteen buffers. At that point, the Transport API's output queue will contain thirteen buffers, each with approximately 500 bytes that it can pass to the underlying transport, instead of passing one at a time.

You can configure each individual connection's high-water mark. The application can also invoke data being passed to the underlying transport by using the `rsSlFlush` routine or passing the `RSSL_WRITE_DIRECT_SOCKET_WRITE` flag to `rsSlWrite`.

1. Reducing underlying transport header overhead.

2. As previously noted, the application can invoke data being passed to the transport through the use of `rsSlFlush()` or the `RSSL_WRITE_DIRECT_SOCKET_WRITE` flag.

Note the throughput and latency implications, and balance the use of the high-water mark and `rsslFlush` accordingly:

- In high-throughput situations, it is better to make large writes to achieve higher efficiencies (i.e., in this case use the high-water mark).
- In low-throughput situations, data might linger in Transport API queues for longer periods and thus incur latency (i.e., in this case, use `rsslFlush`).

11.2.5 Direct Socket Write

The `rsslWrite` function can be instructed to attempt to pass data directly to the underlying transport by specifying the `RSSL_WRITE_DIRECT_SOCKET_WRITE` flag when calling `rsslWrite`. This flag causes the Transport API to check its current outbound queue depths:

- If the queues are empty, the Transport API passes data directly to the underlying transport, bypassing all queuing logic and delays.
- If the queues are not empty, the Transport API adds data to the appropriate queue, with queued content being passed to the underlying transport in the appropriate order.³

Using this option can reduce latency,⁴ as the message might not get queued. However this option also reduces throughput and increases CPU usage due to more frequent socket writes. You can offset the loss in throughput by packing buffers, though doing so can increase packing latency.

11.2.6 Nagle's Algorithm

For TCP socket connection types, you can set the underlying transport to use Nagle's Algorithm to combine small content fragments into larger network frames. While this algorithm reduces transport overhead (optimizing bandwidth usage), it also increases latency, especially when sending small messages at lower data rates.

To minimize latency, the Performance Tools use `tcp_nodelay` (an `RsslConnectOptions` parameter), which disables Nagle's Algorithm. However, you can reenable this feature using the `-tcpDelay` option. For further details, refer to each tool's list of available options.

11.2.7 System Send and Receive Buffers

For TCP socket connections, the OS uses system send and receive buffers for exchanging content. When the Transport API flushes data to the underlying transport, it passes through these system buffers. During times of high throughput, the application might provide data faster than the underlying transport can send it. If this happens, the system buffers can fill up, and as a result, the underlying transport refuses to accept data. In this case, the transport accepts new data only after some of its buffered content is sent and acknowledged.

If the user instructs the Transport API to pass queued data to the underlying transport but the OS cannot accept additional content at the time, `rsslWrite` or `rsslFlush` will return a positive return value. A positive value indicates that content has been queued in the Transport API and the application should flush it at a subsequent time. However, this state is not considered a failure condition, and the Transport API still has the data in its buffers. In this situation, the write file descriptor of the connection can be added to the I/O notifier of choice, which notifies the application when it can pass additional content to the OS.

You can configure the system's send and receive buffer sizes in the OS, as detailed in OS-specific documentation. Additionally, the Transport API allows users to configure this via the `RsslConnectOptions`, `RsslBindOptions`, and to dynamically increase or decrease buffer size during runtime using the `rsslIoctl` function call.

3. As determined by the various priorities with which the content was written and the flush strategy you configure.

4. As long as the underlying transport can accept the content.

11.2.8 Transport API Buffering

The Transport API uses various optimization techniques for efficient input and output of content, many revolving around pre-allocated buffers which minimize memory creation and destruction. Pre-allocated buffers queue outbound data as well as read large byte-streams from underlying transports.

When a connection is established, the maximum size buffer is negotiated, allowing the Transport API to create input and output buffers that work well with respect to that connection. Because input and output strategies have different challenges, these pre-allocated buffer pools are handled differently depending on whether they are input or output buffers.

11.2.8.1 Input Buffering

The Transport API input buffer is created as one large continuous block of memory, controlled by **numInputBuffers** (configured via **RsslConnectOptions** or **RsslBindOptions**). The number of bytes created in the input buffer is determined by the configured value multiplied by the negotiated **maxFragmentSize**. Having one large block of memory allows **rsslRead** to get as many bytes from a single call to the underlying transport as possible. When the input buffer holds data, the Transport API determines message boundaries and returns a single message to the user. As the application makes subsequent **rsslRead** calls, additional messages are dispatched from the input buffer. After fully processing the input buffer, the Transport API goes back to the underlying transport to again fill the input buffer.

The intent is to have the Transport API read only when needed and to read as much as possible. The amount of data the Transport API actually reads from the network depends on the number of input buffers and the amount of data that the OS has available at that time.

11.2.8.2 Output Buffering

Output buffering is handled differently from input buffering. Because each buffer can be written as a different priority, a continuous block of memory will not work. The Transport API creates the configured number of buffers, treating each buffer as a separate entity. Such a division allows the use of multiple buffers simultaneously, as well as allowing buffers to co-exist in different priority-based output queues.

You should configure the number of output buffers according to the application's expected output load. The **guaranteedOutputBuffer** setting controls the number of output buffers available exclusively to that channel, where all of these buffers are created up-front. When the channel runs on a server application, you can also configure the **maxOutputBuffers**, allowing the application to use buffers from the server's shared buffer pool. The shared pool is grown on-demand, but allows for connections under heavy load to temporarily grab a buffer for use. When returned to the shared pool, another channel can then use this same buffer. This enables the server to balance the memory usage of pre-allocated guaranteed output buffers with the traffic spike tolerance of a dynamic shared buffer pool.

Increasing the number of output buffers can improve performance when sending high volumes. An application should be aware of trade-offs of using too much memory and thus potentially slowing the process. If the receiving process cannot keep up with the send rate, a condition can develop for the sender where all output buffers are in use, waiting to be transmitted. The number of output buffers can be dynamically grown using **rsslIoctl** or, more commonly, the connection is terminated as a result.

11.2.8.3 Fragmentation

The negotiated maximum buffer size is the maximum size that the application will send in a single buffer. In cases where an application uses **rsslGetBuffer** to request a buffer with a size larger than the maximum, the requested size will be returned to the user. When the content passes to **rsslWrite**, the Transport API fragments the content on behalf of the application, breaking apart larger content into individual buffers whose individual sizes do not exceed the agreed upon maximum. On the receiving side, the Transport API reassembles the fragments back into a single buffer containing all relevant content.

This transport level fragmentation incurs multiple copies and potential memory allocations. To avoid such overhead, applications should ensure that the maximum buffer size (**RsslBindOptions.maxFragmentSize**) is large enough for commonly sent messages to fit into a single buffer.

11.2.9 Compression

The Transport API supports the use of data compression. Generally, compressing data reduces the amount of data passed to the underlying transport. But compression has some drawbacks to consider:

- Compression requires additional processing.⁵
- Compression copies data: as the user-provided buffer is read by the compression algorithm, output data is compressed into a different buffer. As a result, compression will generally require more buffers from the Transport API's buffer pool.

`upacTransportPerf` exposes the option to enable compression using `-compressionType` and `-compressionLevel`.

11.3 Encoder and Decoder Best Practices

11.3.1 Single-Pass Encoding

Transport API Performance Tools encode data so as to minimize copying. Thus, the application encoding process begins by starting with the top-level container and working down in a linear fashion.

For example, when encoding a Market Price message, the message header is encoded (`rsslEncodeMsgInit`), followed by the field list payload (`rsslEncodeFieldListInit ... rsslEncodeFieldListComplete`). After the payload is encoded, message encoding is completed (`rsslEncodeMsgComplete`).

Encoding the field list prior to the message header would require it to be encoded into a temporary buffer which would then be passed to the message encoder (`rsslEncodeMsg`). This approach would incur multiple buffer allocations and copies to complete encoding.

11.3.2 Clear Functions

The Transport API provides clear functions for its structures (e.g., `rsslClearEncodeIterator`) as well as static initializers (e.g., `RSSL_INIT_ENCODE_ITERATOR`). These functions are tuned to avoid initializing unnecessary structure members and optimize structure use and reuse. In general, Thomson Reuters recommends that you use the clear functions over static initializers, because clear functions are more efficient.

11.4 Other Practices: CPU Binding

Although the OS tries to balance its load intelligently across multi-core processors, you can improve performance by locking the threads of a process to specific cores and CPUs. This lessens the likelihood of switching a thread from one core to another, which impacts processing time and invalidates the cache that the thread has filled.

Each Transport API Performance Tool allows CPU binding through its respective `-threads` option. Thomson Reuters recommends that you test different bindings for each tool to see which works best on each system.

5. Overhead will vary based on the type of compression used and the level of compression applied.

Appendix A Troubleshooting

A.1 Can't Connect

There are many reasons why a consumer or provider might not be able to connect. Several common ones are listed below:

- Check the consumer's and provider's **serviceName** parameters. These must match. The consumer will wait until the service is available and accepting requests.
- Check the ADH (**adhmon**) and ADS (**adsmon**) to see whether the desired service is up.
- Check the ADH's configuration to make sure that the provider's host is listed in the **hostList** configuration setting.
- Check that the provider is listening on the correct TCP Port.
- Check that the consumer is connecting to the correct **hostName** and TCP Port.
- In direct-connect mode, start the provider first, then start the consumer. Starting the consumer first results in a connection timeout, which creates a (by default) 15 second delay until the client retries the connection attempt.
- When connecting through TREP, check that the desired service is up on both the ADH and ADS before starting the consumer (or wait the appropriate amount of time.) Starting the consumer too quickly results in a connection retry after (by default) 15 seconds.

A.2 Not Achieving Steady State

There are several reasons why a consumer might not reach a steady state:

- The **steadyStateTime** value may be too small. When publishing in latency mode or at high update rates, providers will take longer to process image requests. For example, if **steadyStateTime** is set to **30s** but the provider can publish only 2,500 images per second, the consumer times out before it receives its 100,000 images.
- The provider might be overloaded. If the provider is publishing at or near 100% CPU for its configured update rate, it will be either unable or barely able to service incoming image requests, which causes images to trickle back to the consumer.
- The consumer might be overloaded.
- If using a non-interactive provider application, the provider and consumer watch lists might not match, resulting in the consumer application requesting items that never appear in the ADH cache.

A.3 Consumer Tops Out but Not at 100% CPU

In some cases, when connecting to TREP, the consumer appears to be overloaded even though no thread is using the maximum CPU. Such a situation might be a symptom of a bottleneck on the ADS, which can be resolved by increasing the size of the **guaranteedOutputBuffers** and **maxOutputBuffers** to 5,000 in **rmads.cnf**:

```
[...]  
*ads*maxOutputBuffers : 5000  
*ads*guaranteedOutputBuffers : 5000  
[...]
```

Figure 28. ADS rmads.cnf

While this may increase the overall throughput, it can also increase message latency.

A.4 Initial Latencies Are High

- Initial latencies during startup and immediately following the transition to steady state might be high. At high update rates, the system processes its entire overhead for updates plus all refresh traffic, resulting in an increased workload and higher latency. It can take several seconds for the system to “settle” following the transition to steady state. Increasing the provider's output buffers might help.
- Some older systems may not have a reliable timer that the tools can use to measure latencies. Try using a newer system and/or newer OS.

A.5 Latency values Are Very High

- Run the applications on the same machine.
- Use a reliable clock to gather timestamp information.
- Perform appropriate system-wide tuning.
- Consider packing messages into the same buffer. It is possible that the connection type cannot sustain the data rate when sent as individual messages.
-

© 2016 - 2018 Thomson Reuters. All rights reserved.

Republication or redistribution of Thomson Reuters content, including by framing or similar means, is prohibited without the prior written consent of Thomson Reuters. 'Thomson Reuters' and the Thomson Reuters logo are registered trademarks and trademarks of Thomson Reuters and its affiliated companies.

Any third party names or marks are the trademarks or registered trademarks of the relevant third party.

Document ID: ETAC320PETOO.180

Date of issue: 27 April 2018



THOMSON REUTERS